



a place of mind

THE UNIVERSITY OF BRITISH COLUMBIA

CPSC 213: Introduction to Computer Systems

Unit 1a: Numbers and Memory

All slides adapted from materials by Mike Feeley, Jonatan Schroeder, Robert Xiao, and Jordon Johnson


Overview

- Reading
 - Companion: 2.2.2
 - Textbook: 2.1 – 2.3
- Learning objectives:
 - know the number of bits in a byte and the number of bytes in a short, long, and quad
 - determine whether an address is aligned to a given size
 - translate between integers and values stored in memory for both big- and little-endian machines
 - evaluate and write Java expressions using bitwise operators (&, |, <<, >>, and >>>)
 - determine when sign extension is unwanted and eliminate it in Java
 - evaluate and write C expressions that include type casting and the addressing operators (& and *)
 - translate integer values by hand (no calculator) between binary and hexadecimal, subtract hexadecimal numbers, and convert small numbers between binary and decimal

A simple computing machine

- Memory

- stores data encoded as bits
- program instructions and state (variables, objects, etc.)



- CPU

- reads instruction and data from memory
- performs specified computation and writes result back to memory

- Example




- $C = A + B$
- memory stores: add instruction, and variables A, B, C
- CPU: reads instruction and values of A and B, adds values, and writes result to C

Memory

Abstract view

Memory is a big bag of BYTEs

an ADDRESS names each byte
(0 .. a big number)



Memory

- Naming

- unit of addressing is byte (8 bits)
- every byte of memory has an unique address
- some machines have 32-bit memory addresses, some have 64
 - our machine for this class will use 32-bit addresses

- Access

- many things are too big to fit in a single byte
 - e.g. unsigned numbers > 255 , signed numbers < -128 or > 127 , most instructions, etc.
- CPU accesses memory in contiguous, power-of-two-sized chunks of bytes
- address of a chunk is address of its first byte

Memory

Integer data types by size				
# bytes	# bits	C	Java	ASM
1	8	char	byte	b byte
2	16	short	short	w word
4	32	int	int	l long
8	64	long	long	q quad

We will use only 32-bit integers

Numbers and representation

01101101001011110100110010100111 → ?

- Sometimes we are interested in the integer value of a chunk of bytes
 - base-10, **decimal**, is commonly used to represent this number (our "normal" number system)
 - we need to **convert** from binary to decimal to get this value
- Sometimes we are more interested in the bits themselves
 - In such cases the decimal value isn't particularly important
 - For example, consider **memory addresses**
 - big numbers that name power-of-two-sized things
 - we do not usually care what the base-10 value of an address is
 - we can use a power-of-two sized way to identify addresses

Numbers and representation

Identification using powers of 2

- We might use base-2, **binary**
 - a small 256-byte memory has addresses 0_2 to 11111111_2
 - may be represented as **0b11111111**
 - becomes tedious and hard to read as addresses get larger
- Once upon a time we used base-8, **octal**
 - 64KB memory addresses go up to $1111111111111111_2 = 177777_8$
 - may be represented as **0177777**
 - also got tedious as address sizes increased
- Now we use base-16, **hexadecimal**
 - 4GB memory addresses up to $11111111111111111111111111111111_2 = \text{ffffffff}_{16}$
 - we can write this as **0xffffffff**

Binary ↔ hexadecimal

01101010010101010000111010100011

- 4 bits in a hex "digit" (hexit)

0110 1010 0101 0101 0000 1110 1010 0011

- Consider one hexit at a time

6 a 5 5 0 e a 3

0x6a550ea3

- A byte (8 bits) is just 2 hexits ($2^8 = 16^2$):

0x6a550ea3 → 0x6a 0x55 0x0e 0xa3

32 bits (4 bytes) 0x35D
end with 0

bin	hex
0000	0
0001	1
0010	2
0011	3
0100	4
0101	5
0110	6
0111	7
1000	8
1001	9
1010	A
1011	B
1100	C
1101	D
1110	E
1111	F

Which of the following statements is true?

- ☒ A. The Java constants `16` and `0x10` are exactly the same integer
- ☐ B. The Java constants `16` and `0x10` are different integers
- ☐ C. Neither of the statements above is always true
- ☐ D. I don't know
- ☐ E. I'm just choosing E for participation credit

Hexadecimal operations

$$\begin{array}{r} 18_{10} \\ - 10_{10} \\ \hline 08_{10} \end{array}$$


- We use hexadecimal for memory addresses
 - base-10 value is unimportant
- Addition/subtraction in hex
 - can convert both to decimal, add/subtract, and convert back to hex (tedious)
 - can calculate in hex directly
 - alternative for subtraction: use addition with two's complement

$$\begin{array}{r} 0 \\ 0x12 \\ - 0x0A \\ \hline 0x08 \end{array}$$



- Carry for addition, and borrow for subtraction work the same way as with decimal, on a number line with 16 values

decimal




$$\begin{array}{r} 5 \\ + 6 \\ \hline 11 \\ \hline \end{array}$$

$10^2 \quad 10^1 \quad 10^0$

$$\begin{array}{r} 23 \\ - 17 \\ \hline 6 \end{array}$$

hex



$$\begin{array}{r} 9_{16} \\ + 9_{16} \\ \hline 12_{16} \\ \hline \end{array}$$

$16^2 \quad 16^1 \quad 16^0$

$$\begin{array}{r} 33_{16} \\ - 17_{16} \\ \hline 1C_{16} \end{array}$$

iClicker 1a.2

- Object A is at address `0x10d4`, and object B is at `0x1110`. They are stored contiguously in memory (i.e. they are adjacent to each other without any gap). How big is A?



- A. 16 bytes
- B. 48 bytes
- ☒ C. 60 bytes
- D. 80 bytes
- E. Not enough information to tell

$$\begin{array}{r} 0x1110 \\ - 0x10d4 \\ \hline 0x003c \end{array}$$

↑ ↑
 3×16 12×16

$48 + 12 = 60 \text{ bytes}$

Making integers from bytes


- Each memory address holds 1 byte (8 bits)
- First architectural decision:
 - assembling multiple bytes of memory into integers
- Consider a 32-bit integer (`int` in Java or C)
 - this must occupy 4 bytes in memory
 - If memory address is `i`, then we also need the bytes at `i+1`, `i+2`, and `i+3`
 - Example: if address is `0x2014`, then the integer occupies:
 - `0x2014`, `0x2015`, `0x2016`, `0x2017`
 - What do each of these bytes represent?

4 consecutive bytes
occupied by integer

...
<code>i-1</code>
<code>i</code>
<code>i+1</code>
<code>i+2</code>
<code>i+3</code>
<code>i+4</code>
...

Big-endian

- Integer needs to be chopped into 4 bytes and stored in some order
- We can start at the "big end"
 - The first byte stored is the **most significant** byte
 - Used in old IBM servers, network connections

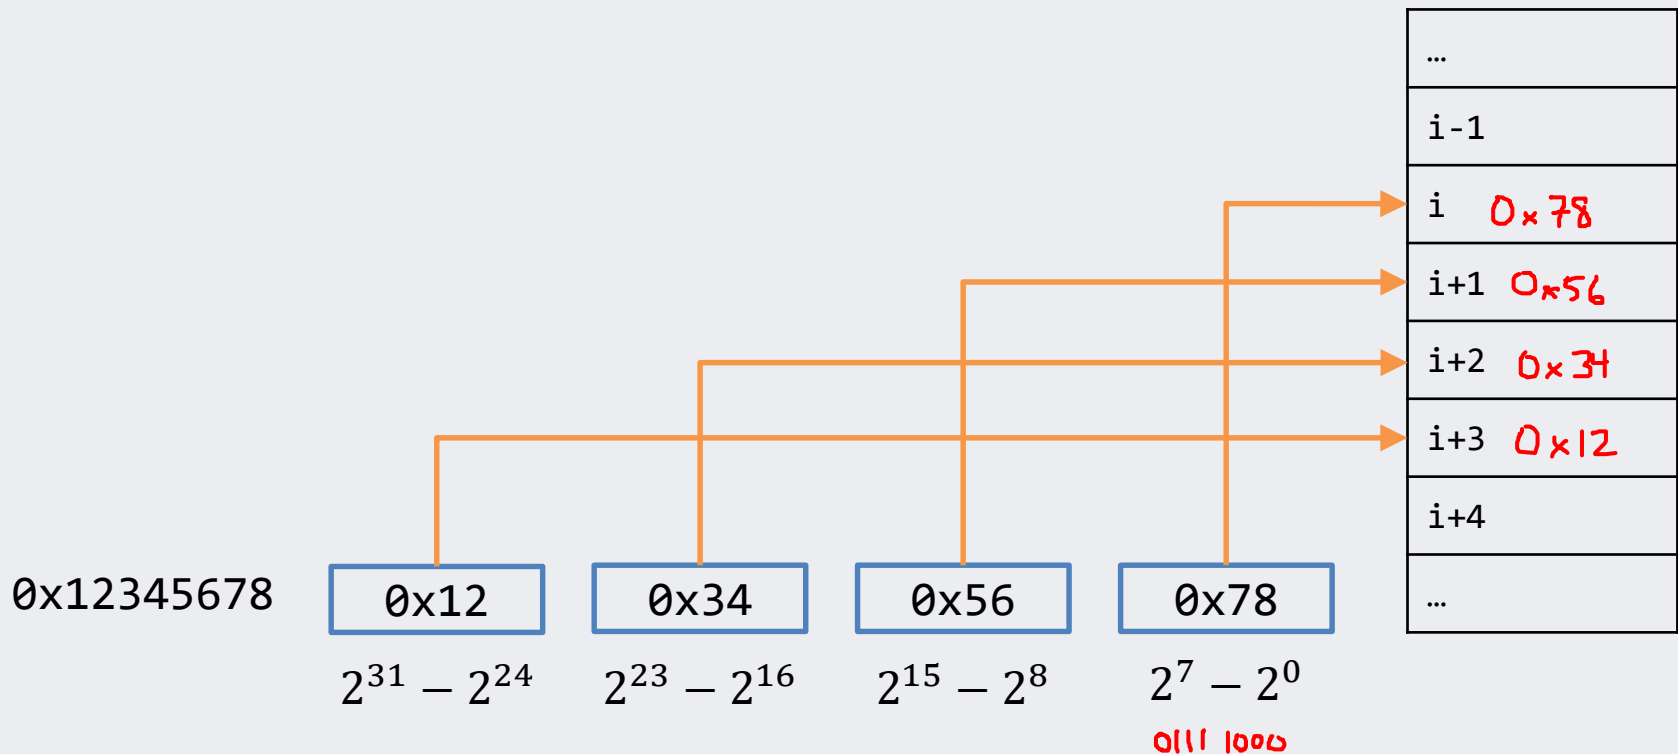


1₁₀

0x00000001

Little-endian

- Integer needs to be chopped into 4 bytes and stored in some order
- We can start at the "little end"
 - The first byte stored is the **least significant** byte
 - Used in most Intel-based systems



iClicker 1a.3

aaa

- What is the Little-endian 4-byte integer value at address 0x4?

- A. 0xc1406b37
- B. 0x1c04b673**
- C. 0x73b6041c
- D. 0x376b40c1
- E. 0x739a8732

Address	Value
0x0:	0xfe
0x1:	0x32
0x2:	0x87
0x3:	0x9a
0x4:	0x73
0x5:	0xb6
0x6:	0x04
0x7:	0x1c

Least significant

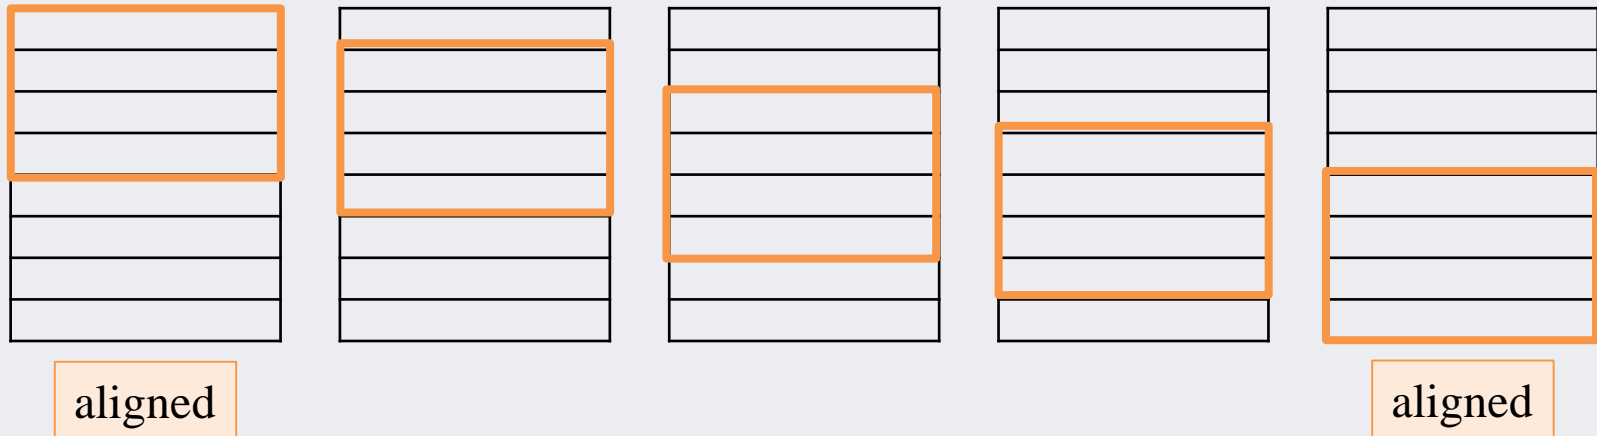
most significant

0 x 1c 04 b6 73
most sig. least sig.

Address alignment

Can we just put data anywhere?



- We could place a 4-byte integer at (almost) any address



- However, forcing addresses to be **aligned** is better for hardware
 - Address whose numeric value is a **multiple** of the object **size**
 - It depends on the object; it gets slightly more complicated with **arrays** and **structs**
- Aligned addresses are better – smaller things fit inside larger things
 - Two **shorts** fit inside an **int**, etc.
 - This is significant for **arrays**

Some CPUs don't support misaligned addresses!

Power-of-two alignment




- Address computation in arrays is achieved by shifting bits
 - e.g. accessing element i , where each element occupies 2^j bits:

$$\&a[i] == \&a[0] + i \cdot (\text{size} == 2^j) == \&a[0] + i \ll j$$
- Memory implementation detail (simplified)
 - memory is actually organized internally into larger chunks called *blocks*
 - suppose a block is 16 bytes
 - then every memory access, internally requires accessing one of these blocks
 - you will see in 313 that this relates to memory caches (cool!)

Alignment and memory access

- Anyway... a CPU memory access looks like this:
 - Read/write **N bytes** starting at **address A**
- The memory (with 16-byte blocks) converts this to:
 - Read/write **N bytes** starting at the **Oth** byte of block **B**
 - **O** is the block offset and **B** is the block number
 - Blocks are numbered, such that block 0 contains addresses 0..15
- The block number and offset are obtained from the calculation:
 $(B, O) = f(A)$




Alignment and memory access

Blocks and offsets

(from previous slide):

- convert a request of **N bytes** starting at **address A**, into
- **N bytes** from the **Oth** byte of block **B**
- How is this simplified if:
 - **block size** (in bytes) is a power of 2, and
 - **N** (object size) is a power of 2, and
 - **A** (absolute address) is aligned to **N**?



iClicker 1a.4

• Which of the following statement(s) is/are true?

- A. The address 6 (110_2)[✓] is aligned for addressing a **short** *2 bytes*
- B. The address 6 (110_2)[✗] is aligned for addressing an **int** (i.e. 4 bytes)
- C. The address 20 (10100_2)[✓] is aligned for addressing an **int**
- D. The address 20 (10100_2)[✗] is aligned for addressing a **long** (i.e. 8 bytes)
- ☒ E. More than one of the above are true

binary: $\begin{array}{r} _ _ 0 \\ _ 0 0 \\ 0 0 0 \end{array}$ *divisible by* $\begin{array}{r} 2 \\ \dots 4 \\ \dots 8 \end{array}$

- Which of the following statement(s) is/are true?
 - A. The address `0x14` is aligned for addressing a `short`, but not an `int`
 - B. The address `0x14` is aligned for addressing a `short` and `int`, but not a `long`
 - C. The address `0x14` is aligned for addressing `short`, `int`, and `long`
 - D. None of the above
 - E. P-A-R-T-I-C-I-P-A-T-I-O-N!

`0001 01002` aligns with `short` & `int`

Shifting bits

- Shifting multiplies or divides by a power of 2
 - shifting left by b bits is the same as multiplying by 2^b
 - shifting right by b bits is the same as dividing by 2^b

$$\begin{aligned} 0x6 \gg 1 &== 110_2 \gg 1 \\ &== 11_2 \\ &== 0x3 \\ 0x6 \ll 1 &== 110_2 \ll 1 \\ &== 1100_2 \\ &== 0xc \end{aligned}$$

- What about negative numbers?
 - Recall that negative numbers are represented as two's complement
 - $-6 == 0xfa == 11111010_2$
 - $-6 / 2 == -3$, but 11111010_2 shifted right is $01111101_2 == 125$, not -3

Shifting bits

Java vs C $94_{10} \gg 1 == 47_{10}$

- Java has two kinds of right shifts

- SIGNED shift ">>", keeps the most significant bit the same

• e.g. $-6 \gg 1 == \hat{1}111\ 1010_2 \gg 1 == \hat{1}111\ 1101_2 == -3$

- UNSIGNED shift ">>>", shifts and sets most significant bit to zero

• e.g. $-6 \ggg 1 == 1111\ 1010_2 \ggg 1 == 0111\ 1101_2 == 0x7d$

$47_{10} / 4 == 47_{10} \gg 2 == \hat{0}010\ 1111_2 \gg 2 == 0000\ 1011_2 == 11_{10}$

- Java: you choose the type of shift by using >> or >>>
- C: compiler chooses based on variable's declared type

- `unsigned int u; int s;` `unsigned` keyword creates an unsigned value

$55_{10} \ll 1 == 00110111_2 \ll 1 == 01101110_2 == 110_{10}$

What about left shifts?

$110_{10} \ll 1 == 01101110_2 \ll 1 == 11011100_2 ==$

8 bits 2's complement $[-128, 127]$

Extending integers

b: 1111 1010₂

- Extending is: *i* ← *1111 1010₂*
 - increasing the number of bits (or bytes) used to store an integer

```
byte b = -6;  
int i = b;  
out.printf("b: 0x%x %d, i: 0x%x %d\n", b, b, i, i);
```

- what prints? *b: 0xfa -6, i: 0xfffffffffa, -6*

- Signed extension**

- used with signed numbers (everything in Java is signed)
- copies sign bit into upper, empty bits of the extended number

- Zero extension**

- used with unsigned numbers (e.g. in C)
- sets upper, empty bits to 0
- can be forced by **masking** using logical (bitwise) AND operator

b: 1111 1010
sign-extended: 1111 ... --- 1111 1010
mask : 0000 ... --- 1111 1111
AND : 0000 ... --- 1111 1010

```
int u = b & 0xff;  
out.printf("u: 0x%x %d\n", u, u);
```

u: 0xfa 250

Truncating integers

- We can also use fewer bits to represent a value

```
int i = -6;  
byte b = i;  
out.printf("b: 0x%x %d, i: 0x%x %d\n", b, b, i, i);
```

i: 0x ffff fffa
b: 0xfa

- what could go wrong?

- if i is 256, what is b? What if i is 128?

i: 0x 0000 0100

b: 0x00

i: 0x 0000 0080
b: 0x80
-128

- Java warns you if you truncate implicitly

- To hide warning, `cast` explicitly

```
byte b = (byte) i;
```

- What is the value of `i` after this Java statement executes?

```
int i = ((byte) 0x8b) << 16;
```

4 bytes

1 byte

sign-extended, then shifted.

- A. 0x8b
- B. 0x0000008b
- C. 0x008b0000
- ☒ D. 0xff8b0000
- E. None of these
- F. I don't know

0x 00 00 00 8b

0x8b

0xff ff ff 8b

0xff 8b 00 00

iClicker 1a.7

- What is the value of `i` after this Java statement executes?

```
int i = 0xff8b0000 & 0x00ff0000;
```

$\& \ 0x00ff0000$

$00\ 8b\ 0000$

- A. `0xfffff0000`
- B. `0xff8b0000`
- C. `0x008b0000`
- D. None of these
- E. I don't know

$x \wedge T \equiv x$

$x \wedge F \equiv F$

8 b

1000	1011	← x
& 1111	1111	← T
<hr/>		
1000	1011	