



a place of mind

THE UNIVERSITY OF BRITISH COLUMBIA

CPSC 213

Introduction to Computer Systems

Unit 1b: Static Scalars and Arrays

All slides adapted from materials by Mike Feeley, Jonatan Schroeder, Robert Xiao, and Jordon Johnson

Announcements

- Google doc for lecture questions
 - See Piazza for link, section 102
 - https://docs.google.com/document/d/1G6hkekQS7mT9lFpP8AVftYao8vLRujIrRLAvOuX_07w/edit



- Add your question anonymously (at the top)
- Help answer questions too!

Overview

- Reading
 - Companion: 1, 2.1-2.3, 2.4.1-2.4.3
 - Textbook: 3.1-3.2.1
 - Reference (textbook, as needed): 3.1-3.5, 3.8, 3.9.3
- Learning objectives:
 - list the basic components of a simple computer and describe their function
 - describe ALU functionality in terms of inputs and outputs
 - describe the exchange of data between ALU, registers, and memory
 - identify and describe the basic components of a machine instruction
 - outline the steps a RISC machine performs to do arithmetic on numbers stored in memory
 - outline the steps a simple CPU performs when it processes an instruction
 - translate between array-element offsets and indices
 - distinguish static and dynamic computation for access to global scalars and arrays in C
 - translate between C and assembly language for access to global scalars and arrays
 - translate between C and assembly for code that performs simple arithmetic
 - explain the difference between arrays in C and Java

Our approach

- Develop a model of computation
 - that is rooted in what the machine actually does
 - by examining C, bit-by-bit (comparing to Java as we go)
- The processor
 - we will design (and you will implement) a **simple instruction set**
 - based on what we need to compute C programs
 - similar to a real instruction set (MIPS)
- The language
 - we will act as compiler to translate C into machine language/assembly
 - bit by bit, then putting the bits together to do interesting things
 - edit, debug, and run using simulated processor to visualize execution

The CPU


- CPUs execute **instructions**, not C or Java code
- Execution proceeds in **stages**
 - **Fetch**: load the next instruction from memory
 - **Decode**: figure out from the instruction what needs to be done
 - **Execute**: do what the instruction specifies
 - There can be more or fewer stages than this, depending on your model/implementation (CPSC 121, 313)
- These stages are looped over and over again **forever**

CPU instructions

- CPU instructions are very **simple**
 - Read (**load**) a value from memory
 - Write (**store**) a value to memory
 - Add/subtract/AND/OR/etc. two numbers
 - Shift a number
 - Control flow (see these in unit 1d)
- Some of these operations are carried out by an ALU (Arithmetic & Logic Unit)
 - ALU is used in the **execute** stage

Phases of computation

- **Human creation**: design program and describe it in high-level language
- **Compilation**: convert high-level human description into machine-executable text
- **Execution**: a physical machine executes the code text




Static vs dynamic computation

- Execution

- Parameterized by input values **unknown** at compilation
- Producing output values that are **unknowable** at compilation

- Anything the **compiler can compute** is called *static*
- Anything that can **only be discovered during execution** is called *dynamic*



- ^C
✓
- Consider the code below:

```
int a;  
a = 5;
```

The value assigned to variable **a** is:

- ☒ A. static
- ☐ B. dynamic
- ☐ C. neither static nor dynamic
- ☐ D. it depends

- Assume that a variable of type `int` is assigned a value that is read from the user through an input dialog.

This value is:

- A. static
- ☒ B. dynamic
- C. neither static nor dynamic
- D. it depends

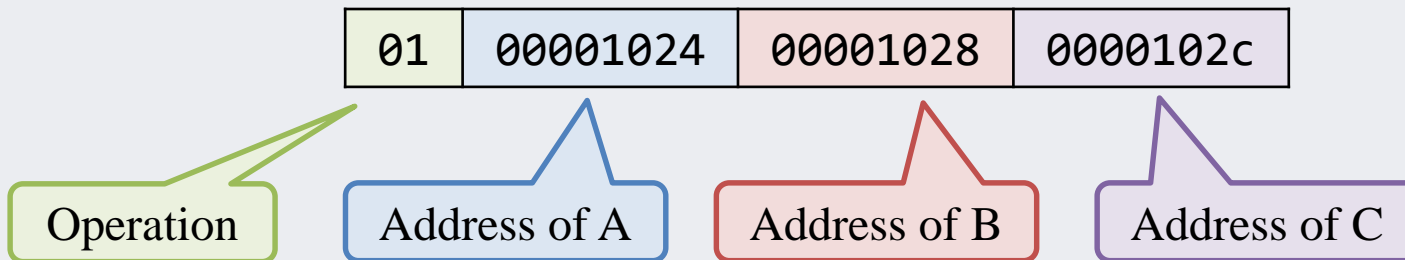
The processor (CPU)

- Implements a set of **instructions**
- Each instruction is implemented using **logic gates**
 - Built from **transistors**: fundamental mechanism of computation
 - (Recall your CPSC 121 labs)
- Instruction design philosophies
 - 213 → **RISC**: *Reduced instruction set computer* fewer and simpler instructions makes processor design simpler
 - **CISC**: having **more types** of instructions (and more complex instructions) allows for shorter/simpler program code and simpler compilers
 - Complete instruction set computer*

First proposed instruction: ADD

- Let's propose an instruction that does: $A \leftarrow B + C$
 - where A , B , and C are stored in memory
- Instruction parameters: addresses of A , B , C
 - each address is 32 bits (modern computers: 64 bits)
- Instruction is encoded as a sequence of bits (stored in memory)
 - **Operation name** (e.g. add)
 - Addresses for A , B , C

This will occupy at least 13 bytes in memory




Improving the ADD instruction

Problems with memory access

- Accessing memory is SLOW
 - ~100 CPU cycles for every memory access
 - goal: fast programs that avoid accessing memory when possible
- Big instructions are costly
 - Memory addresses are big (so instructions that use them are big)
 - Big instructions lead to big programs
 - Reading instructions from memory is slow (fewer caching options)
 - Large instructions use more CPU resources (transfer, storage)

General purpose registers

- Register file
 - Small, fast memory stored in CPU itself
 - roughly **single-cycle** access
- Registers
 - Each register named by a number (e.g., 0–7)
 - some of these may sometimes be used for other purposes
 - Size: architecture's common integer (32 or 64 bits)



Instructions using registers

Improving our ADD instruction


$$A \leftarrow B + C$$

Let's modify our instruction design


01	00001024	00001028	0000102c
----	----------	----------	----------

- Memory instructions handle only memory
 - Load data from memory into a register (slow)
 - Store data from register into memory (slow)

- Other instructions access data in registers
 - small and fast



- To further improve instruction size: share register for one source and the destination



Special purpose registers

- A special-purpose register can only be used for certain purposes
 - physically separate from the general-purpose register file
 - May not be accessible by all instructions (e.g., cannot be used as an argument for an add instruction)
 - May have special meaning or be treated specially by CPU
- Examples:
 - **PC (Program Counter)**: contains address of next instruction to execute
 - **IR (Instruction Register)**: contains the machine instruction that has just been fetched from memory

Instruction Set Architecture (ISA)

- ISA is a **formal interface** to a **processor implementation**
 - defines the instructions the processor implements
 - defines the format of each instruction
- **Types of instructions:**
 - math and logic
 - memory access
 - control transfer: "goto" and conditional "goto"

- Design alternatives:
 - CISC: simplify compiler design (e.g. Intel x86, x86-644)
 - RISC: simplify processor implementation
- Instruction format
 - sequence of bits: **opcode** and **operand** values
 - all represented as **numbers**
- **Assembly language**:
 - symbolic (textual) representation of machine code

Representing Instruction Semantics

↖ what the instruction does

right to left

- RTL: simple, convenient pseudo language to describe semantics
 - easy to read/write, describes machine steps
- Syntax:
 - each line is of the form: LHS ← RHS
 - LHS is a memory or register that receives a value
 - RHS is constant, memory, register, or expression on two registers
 - m[a] is a memory in address a
 - r[i] is a register with number i

Register file and memory are treated as arrays

RTL examples

- Register 0 receives $0x2000$

- $r[0] \leftarrow 0x2000$

- Register 1 receives memory whose address is in register 0

- $r[1] \leftarrow m[r[0]]$

- Register 2 is increased by the value in register 1:

- $r[2] \leftarrow r[2] + r[1]$

r0	0x2000
r1	410
r2	5 415
r3	
r4	
r5	
r6	
r7	

	memory
	...
0x2000	
0x2001	
0x2002	410
0x2003	
0x2004	
	...

RTL example

- Assume the value 7 is stored at address 0x1234.
- What is the result of the following sequence of instructions?

- $r[0] \leftarrow 10$
- $r[1] \leftarrow 0x1234$
- $r[2] \leftarrow m[r[1]]$
- $r[2] \leftarrow r[0] + r[2]$
- $m[r[1]] \leftarrow r[2]$

r0	10
r1	0x1234
r2	7 17
r3	
r4	
r5	
r6	
r7	

	memory
	...
0x1234	7 17
0x1238	
	...

Variables

- Variables are **named storage locations** for values
- Features:
 - Name
 - Type/size
 - Scope
 - Lifetime
 - Memory location (address)
 - Value
- Which of these are **static**? Which are **dynamic**?
 - Which are **determined** or **can change** while the program is running?

Static variables, built-in types

- In **Java**:

- static data members are allocated to a class, not an object
- they can store built-in scalar types or references to arrays or objects (later)

Java:

```
public class Foo {  
    static int a;  
    static int[] b; // array not static, ignore for now  
    public void foo() {  
        a = 0;  
        b[a] = a;  
    }  
}
```

C:

```
int a;  
int b[10];  
void foo() {  
    a = 0;  
    b[a] = a;  
}
```

- In **C**:

- global variables and any other variable declared static
- they can be static scalars, arrays or structs or pointers (later)

Static variable allocation

```
int a;  
int b[10];  
  
void foo() {  
    a = 0;  
    b[a] = a;  
}
```

Static memory layout

0x1000: value of a
...
0x2000: value of b[0]
0x2004: value of b[1]
...
0x2024: value of b[9]

- Allocation is
 - assigning a memory location (i.e. an address) to a variable
- Static vs dynamic computation
 - global/static variables can exist before program starts and live until after it finishes
 - compiler allocates variables, giving them a constant address
 - no dynamic computation is required to allocate; they just exist
 - compiler tracks free space during compilation – it is in complete control of program's memory

- Assume `a` is a global variable in C. When is space for `a` allocated? In other words, when is its address determined?
- ☒ A. The compiler assigns the address when it compiles the program
 - B. The compiler calls the memory to allocate `a` when it compiles the program
 - C. The compiler generates code to allocate `a` before the program starts running
 - D. The program locates available space for `a` before it starts running
 - E. The program locates available space as soon as the variable is used for the first time

Static variable access

Scalars

```
int a;  
int b[10];  
  
void foo() {  
    a = 0;  
    b[a] = a;  
}
```

Static memory layout

0x1000: value of a
...
0x2000: value of b[0]
0x2004: value of b[1]
...
0x2024: value of b[9]

- Key observation:
 - Addresses of a, b[0], b[1], ... are constants known to the compiler
- Use RTL to specify instructions needed for `a = 0;`

$m[0x1000] \leftarrow 0$

- Now generalize – what about, for example,

- `a = x + y;`? *suppose x, y are static*
- `c = a;`?

$m[0x1000] \leftarrow m[\text{static address of } x] + m[\text{static address of } y]$

Static variable access

Static arrays

```
int a;  
int b[10];  
  
void foo() {  
    a = 0;  
    b[a] = a;  
}
```

Static memory layout

0x1000: value of a

...

0x2000: value of b[0]

0x2004: value of b[1]

...

0x2024: value of b[9]

$0x2000 + 9 * 4 \text{ bytes}$

$0x2000 + 36_{10}$

↓

$2 * 16^1 + 4 * 16^0$

$0x2000 + 0x24$

- Key observation:
 - compiler does not know address of `b[a]`
 - unless it knows the value of `a` statically, which it could by looking at `a = 0;`, but not in general
- Array access is computed from a base address and **index**
 - address of element = **base** + **offset**; **offset** = **index** × **element size**
 - The base address (0x2000) and element size (4) are static; the index is dynamic

iClicker 1b.4

- What is a proper RTL instruction for $b[a] = a$;

Assume that the compiler does not know the current value of a .

$0x1000$: address of a
 $m[0x1000]$: value of a

Static memory layout

0x1000: value of a
...
0x2000: value of $b[0]$
0x2004: value of $b[1]$
...
0x2024: value of $b[9]$

- A. $m[0x2000+m[0x1000]] \leftarrow m[0x1000]$
- ☒ B. $m[0x2000+4*m[0x1000]] \leftarrow m[0x1000]$
- C. $m[0x2000+m[0x1000]] \leftarrow 4*m[0x1000]$
- D. $m[0x2000]+4*m[0x1000] \leftarrow m[0x1000]$
- E. $m[0x2000]+m[0x1000] \leftarrow m[0x1000]$

ISA design

Choosing the instructions

$a = 0;$

$m[0x1000] \leftarrow 0x0$

$b[a] = a;$

$m[0x2000 + 4*m[0x1000]] \leftarrow m[0x1000]$

- ISA design goals
 - minimize # of memory instructions in ISA
 - at most 1 memory access per instruction
 - minimize # of total instructions in ISA
 - minimize instruction size

$b[a-1] = a;$
 $r[0] \leftarrow 0x1000$
 $r[1] \leftarrow m[r[0]]$
 $r[2] \leftarrow 0x2000$
load $r[3]$ with $a-1$
 $m[r[2] + 4*r[3]] \leftarrow r[1]$

$r[0] \leftarrow 0x0$
 $r[1] \leftarrow 0x1000$ ^{address for a} (1)
 $m[r[1]] \leftarrow r[0]$ (2)

$r[0] \leftarrow 0x1000$
 $r[1] \leftarrow m[r[0]]$ (3)
 $r[2] \leftarrow 0x2000$
 $m[r[2] + 4*r[1]] \leftarrow r[1]$ (4)

and we also have a "load" version of (4) (5)

SM213 ISA

The first five instructions so far

what it does

	Name	Semantics	Assembly	Machine
①	load immediate	$r[d] \leftarrow v$	ld \$v, rd	0d-- vvvvvvvv 6 bytes
③	load base + offset	$r[d] \leftarrow m[r[s]]$	ld (rs), rd	1?sd 2 bytes
⑤	load indexed	$r[d] \leftarrow m[r[s]+4*r[i]]$	ld (rs, ri, 4), rd	2sid
②	store base + offset	$m[r[d]] \leftarrow r[s]$	st rs, (rd)	3s?d
④	store indexed	$m[r[d]+4*r[i]] \leftarrow r[s]$	st rs, (rd, ri, 4)	4sdi

→ incomplete

Translating the code

- To RTL

```
int a;  
int b[10];  
  
void foo() {  
    a = 0;  
    b[a] = a;  
}
```



```
r[0]    ← 0  
r[1]    ← 0x1000  
m[r[1]] ← r[0]  
  
r[2]    ← m[r[1]]  
r[3]    ← 0x2000  
m[r[3]+4*r[2]] ← r[2]
```

- To assembly

```
r[0]    ← 0  
r[1]    ← 0x1000  
m[r[1]] ← r[0]  
  
r[2]    ← m[r[1]]  
r[3]    ← 0x2000  
m[r[3]+4*r[2]] ← r[2]
```



```
ld $0, r0  
ld $0x1000, r1  
st r0, (r1)  
  
ld (r1), r2  
ld $0x2000, r3  
st r2, (r3, r2, 4)
```

Translating the code

If a human wrote the assembly code

- list static allocations, use labels for addresses, add comments

```
→ .pos 0x100
ld $0, r0    # r0 = 0
ld $a, r1    # r1 = address of a
st r0, (r1)  # a = 0

ld (r1), r2  # r2 = a
ld $b, r3    # r3 = address of b
st r2, (r3, r2, 4) # b[a] = a

→ .pos 0x1000
→ a:
→ .long 0 # the variable a

.pos 0x2000
→ b:
.long 0 # the variable b[0]
.long 0 # the variable b[1]
...    # need to write out the rest
.long 0 # the variable b[9]
```

• ——— preprocessor directives (not translated to machine code)

• labels - symbolic addresses

• a number

0x1000	0	a
...		
0x2000	0	b
0x2004	0	
0x2008	0	
0x200c	0	
0x2010	0	
0x2014	0	
0x2018	0	
0x201c	0	
0x2020	0	
0x2024	0	
...		

iClicker 1b.5

```
int i;
int a[10];

a[2] = a[i];
```

Name	Semantics	Assembly
load immediate	$r[d] \leftarrow v$	ld \$v, rd
load base + offset	$r[d] \leftarrow m[r[s]]$	ld (rs), rd
load indexed	$r[d] \leftarrow m[r[s]+4*r[i]]$	ld (rs, ri, 4), rd
store base + offset	$m[r[d]] \leftarrow r[s]$	st rs, (rd)
store indexed	$m[r[d]+4*r[i]] \leftarrow r[s]$	st rs, (rd, ri, 4)

• Which correctly implements `a[2] = a[i];`?

A. `st ($a, $i, 4), ($a, 2, 4)`

B. `ld ($a, $i, 4), r0`
`st r0, ($a, 2, 4)`

C. `ld $a, r0`
`ld $i, r1` *ld (r1), r1*
`ld (r0, r1, 4), r2`
`st r2, 8(r0)`

D. `ld $a, r0`
`ld $i, r1`
`ld (r0, r1, 4), r2`
`st r2, (r0, 2, 4)`

E. None of these / I don't know

The Simple Machine (SM213) ISA

- Architecture
 - Register file Eight 32-bit general purpose registers (numbered 0-7)
 - CPU one cycle per instruction (fetch + execute)
 - Main Memory byte addressed, Big endian integers
- Instruction format
 - 2- or 6-byte instructions (each character below is a hex digit)
 - `x-01`, `xxsd`, `x0vv`, or `x-sd vvvvvvvv`
 - where:
 - `x` is an opcode (unique identifier for this instruction)
 - `-` means unused
 - `s` and `d` are operand register numbers
 - `vv`, `vvvvvvvv` are immediate/constant values

Machine and assembly syntax

- Machine code

- `[addr:] x-01 [vvvvvvvvv]`

- `addr:` sets starting address for subsequent instructions
- `x-01` hex value of an instruction with opcode x and operands 0 and 1
- `vvvvvvvvv` hex value of optional extended value

- Assembly code

- `([label:] [instruction | directive] [# comment] |)*`

- `directive` :: `(.pos number) | (.long number)`
- `instruction` :: `opcode operand+`
- `operand` :: `$literal | reg | offset(reg) | (reg, reg, 4)`
- `reg` :: `r0..7`
- `literal` :: `number`
- `offset` :: `number`
- `number` :: `decimal | 0xhex`

Memory access instructions

Load and Store with different **addressing modes**

Name	Semantics	Assembly	Machine
load immediate	$r[d] \leftarrow v$	ld \$v, rd	0d-- vvvvvvvv
load base + offset	$r[d] \leftarrow m[r[s] + (o = 4 * p)]$	ld o(rs), rd	1psd
load indexed	$r[d] \leftarrow m[r[s] + 4 * r[i]]$	ld (rs, ri, 4), rd	2sid
store base + offset	$m[r[d] + (o = 4 * p)] \leftarrow r[s]$	st rs, o(rd)	3spd
store indexed	$m[r[d] + 4 * r[i]] \leftarrow r[s]$	st rs, (rd, ri, 4)	4sdi

- We have specified 4 **addressing modes** for operands
 - immediate constant value stored as part of instruction
 - register operand is register number; register stores value
 - base+offset operand is register number;
register stores memory address of value (+ offset = $p * 4$)
 - indexed two register-number operands;
store base memory address and value of index

Basic arithmetic, shifting, NOP, and halt

- Arithmetic

Name	Semantics	Assembly	Machine
register move	$r[d] \leftarrow r[s]$	mov r_s, rd	60sd
add	$r[d] \leftarrow r[d] + r[s]$	add r_s, rd	61sd
and	$r[d] \leftarrow r[d] \& r[s]$	and r_s, rd	62sd
inc	$r[d] \leftarrow r[d] + 1$	inc rd	63-d
inc address	$r[d] \leftarrow r[d] + 4$	inca rd	64-d
dec	$r[d] \leftarrow r[d] - 1$	dec rd	65-d
dec address	$r[d] \leftarrow r[d] - 4$	deca rd	66-d
not	$r[d] \leftarrow \sim r[d]$	not rd	67-d

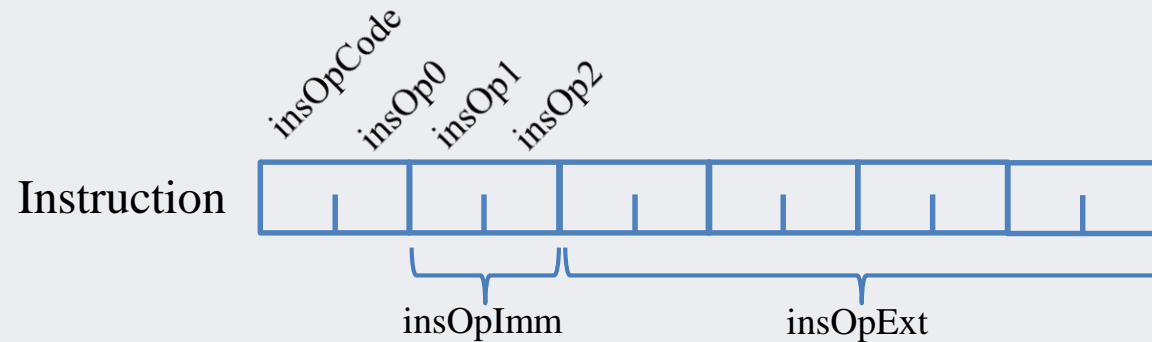
- Logical (shifting), NOP, and halt

Name	Semantics	Assembly	Machine
shift left	$r[d] \leftarrow r[d] \ll v=s$	shl $\$v, rd$	7dss (ss>0)
shift right	$r[d] \leftarrow r[d] \gg v=s$	shr $\$v, rd$	7dss (ss<0)
halt	<i>halt machine</i>	halt	F0--
nop	<i>do nothing</i>	nop	FF--

The SM213 CPU implementation


- Internal state

- PC address of *next* instruction to fetch
- Instruction the value of the current instruction



Reg	Value
PC:	0000010e
Instruction:	3001 00000000
Ins Op Code:	3
Ins Op 0:	0
Ins Op 1:	0
Ins Op 2:	1
Ins Op Imm:	01
Ins Op Ext:	00000000

- Cycle stages



- Fetch

- read instruction at PC, determine size, separate components, set next sequential PC

- Execute

- read internal state, perform specified computation (insOpCode), update internal state
- read and update may be to memory as well

The SM213 CPU implementation

Java syntax

- Internal registers

- `insOp0, insOp1, insOp2, insOpImm, insOpExt, pc`
- `get()`
- `set(value)`

name of a register

```
int i = insOp1.get();  
insOp1.set(i);
```

- General purpose registers

- `reg.get(registerNumber)`
- `reg.set(registerNumber, value)`

name of a register

```
int i = reg.get(3); // i <- r[3]  
reg.set(3, i);     // r[3] <- i
```

- Main Memory

- `mem.readInteger(address)`
- `mem.writeInteger(address, value)`

```
int i = mem.readInteger(0x1000); // i <- m[0x1000]  
mem.writeInteger(0x1000, i);    // m[0x1000] <- i
```

Global dynamic arrays

- Java
 - array variable stores reference to array allocated dynamically with **new** statement

```
public class Foo {  
    static int a;  
    static int b[];  
  
    void foo() {  
        b = new int[10]  
        b[a] = a;  
    }  
}
```

- C
 - array variables can store static arrays, or
 - pointers to arrays allocated dynamically with call to **malloc** library procedure

static {

```
int a;  
int* b; ← will (eventually) hold the address of the dynamic array
```

dynamic array

```
void foo() {  
    b = malloc( 10 * sizeof(int) );  
    b[a] = a;  
}
```

static array

```
int a;  
int b_data[10];  
int* b = &b_data[0];
```


Static vs dynamic arrays

- Declared and allocated differently, but accessed the same

```
int a;  
int b[10];  
  
void foo() {  
    b[a] = a;  
}
```

```
int a;  
int* b; pointer  
  
void foo() {  
    b = malloc( 10 * sizeof(int) );  
    b[a] = a;  
}
```

- Allocation

- for static arrays, the compiler allocates the array
- for dynamic arrays, the compiler allocates a pointer

a: 0x1000: value of a

b:

0x2000: value of b[0]
0x2004: value of b[1]
...
0x2024: value of b[9]


a: 0x1000: value of b

b:

0x2000: value of b (will contain address of array) <i>(eventually)</i>

Static vs dynamic arrays

- ...then when the program runs
 - the dynamic array is allocated by a call to `malloc`, e.g. at address `0x3000`
 - the (static) variable `b` is set to the memory address of this array



- Generating code to access the array
 - for the dynamic array, the compiler generates an additional load for `b`

`b[a] = a ;`

Static

<code>r[0]</code>	<code>← 0x1000</code>	<i>address of a</i>
<code>r[1]</code>	<code>← m[r[0]]</code>	<i>value of a</i>
<code>r[2]</code>	<code>← 0x2000</code>	<i>address of b</i>
<code>m[r[2]+4*r[1]]</code>	<code>← r[1]</code>	<i><code>b[a] = a ;</code></i>

Dynamic

<code>r[0]</code>	<code>← 0x1000</code>	<i>address of a</i>
<code>r[1]</code>	<code>← m[r[0]]</code>	<i>value of a</i>
<code>r[2]</code>	<code>← 0x2000</code>	<i>address of b (pointer)</i>
<code>r[3]</code>	<code>← m[r[2]]</code>	<i>address of b (array)</i>
<code>m[r[3]+4*r[1]]</code>	<code>← r[1]</code>	<i><code>b[a] = a ;</code></i>

*`int a ;`
`int* b ;`*

C arrays


How they differ from Java arrays

- Terminology
 - use the term *pointer* instead of *reference*; they mean the same thing
- Declaration
 - the type is a pointer to the type of its elements, indicated with a *
- Allocation
 - malloc allocates a block of bytes; no type; no constructor
- Type safety
 - any pointer can be typecast to any other pointer type
- Bounds checking
 - C performs no array bounds checking
 - out-of-bounds access manipulates memory that is not part of the array
 - this is the major source of virus vulnerabilities in the world today



Summary: arrays in Java and C

- In Java and C:
 - an array is a list of items of the same type
 - array elements are named by non-negative integers starting at 0
 - syntax for accessing element `i` of array `b` is `b[i]`
- In Java:
 - variable stores a reference to the array
- In C:
 - variable can store a pointer to the array, or the array itself
 - distinguished by variable type (i.e. pointer variable vs array variable)
 - array element access via pointer can be done by
 - `b[i]`, or
 - `*(b+i)`



Declaring (and dereferencing) C pointers

- Pointer variables declared as:

TYPE* p;

where TYPE is any type (even another pointer type), e.g. `int`

- multiple variables can be declared on a line, but beware

- `int* x, y; // this declares x as int*, and y as int`
- `int *x, *y; // this declares x and y both as int*`

- Even more complicated with in-line initialization, so Geoff recommends keeping one declaration per line

- Can be assigned the address of a variable of type TYPE

```
int a = 5;
int* p = &a;
```


The value of a pointer is a memory address



↖ "address of" a

- Access the value being pointed to by dereferencing with *

```
int a = 5;
int* p = &a;
• *p = 37;
```



Pointer arithmetic

$*(b+i)$



$*(0x100c)$

$int^* p$ $0x1004$ $*(p+2)$


- Purpose:

- an alternative way to access array elements compared to $a[i]$

- Addresses are just numbers, so...

- adding or subtracting an integer index to a pointer (address)
 - results in the address of something else (another pointer of the same type)
 - value of the pointer is offset by $\text{index} \times \text{size of the pointer's referent}$
- e.g. adding 3 to an int^* yields a pointer value 12 bytes larger than the original

$p+3$



- subtracting two pointers of the same type

- results in an integer: the number of referent-type elements between the two pointers
- e.g. $(\&a[7] - \&a[2]) == 5 == (a+7) - (a+2)$

↑
of
elements
in between

a : integer array
 int^*

$a+7$: int^*

$\&a$: int^{**}

Pointer arithmetic

Example

- The following C programs are identical

```
int* a;  
a[7] = 5;
```

```
int* a;  
*(a+7) = 5;
```

- For array access, the compiler would generate this code

```
r[0]    ← a  
r[1]    ← m[r[0]]  
r[2]    ← 7  
r[3]    ← 5  
m[r[1]+4*r[2]] ← r[3]
```

```
ld $a, r0  
ld (r0), r1  
ld $7, r2  
ld $5, r3  
st r3, (r1, r2, 4)
```

- multiplying the index (7) by 4 (size of integer) to compute the array offset

iClicker 1b.6

- Which element of the original 10-element array is modified with the highlighted instructions?

- ☒ A. Element with index 3
- ☐ B. Element with index 6
- ☐ C. Element with index 0
- ☐ D. No element is modified
- ☐ E. More than one element is modified

```
int* c;  
  
void foo() {  
    c = malloc(10 * sizeof(int));  
    c = &c[3];  
    *c = *&c[3];  
}
```

c :

0x1000	0x2000 0x200c	
...		
0x2000		<i>c[0]</i>
0x2004		1
0x2008		2
0x200c	<u>78 - 112</u>	<i>c[3]</i>
0x2010		4
0x2014		5
0x2018	-112	6
0x201c		7
0x2020		8
0x2024		9

c[0]
c[1]
c[2]
c[3]

Previous code in SM213 assembly

```
ld $0x1000, r0 # r0 = address of c
ld (r0), r1     # r1 = c (malloc result)
ld $12, r2      # r2 = 3*sizeof(int)
add r1, r2      # r2 = c+3 = &c[3]
st r2, (r0)     # c = c+3

ld $3, r3       # r3 = 3
ld (r2, r3, 4), r4 # r4 = c[3]
st r4, (r2)     # *c = c[3]
```

```
.pos 0x1000
c: .long 0 # or some data used in simulator to emulate malloc
```

```
int* c;

void foo() {
    c = malloc(10 * sizeof(int));
    c = &c[3];
    *c = *&c[3];
}
```

Summary

Static scalar and array variables

- Static variables
 - the compiler knows the address (memory location) of variable
- Static scalars and arrays
 - the compiler knows the address of the scalar value or array
- Dynamic arrays
 - the compiler does not know the address of the array
- What C does that Java doesn't
 - static arrays
 - arrays can be accessed using pointer dereference operator
 - pointer arithmetic
- What Java does that C doesn't
 - type-safe dynamic allocation
 - automatic array bounds checking