



a place of mind

THE UNIVERSITY OF BRITISH COLUMBIA

CPSC 213

Introduction to Computer Systems

Unit 1c: Instance Variables and Dynamic Allocation

All slides adapted from materials by Mike Feeley, Jonatan Schroeder, Robert Xiao, and Jordon Johnson

Announcements

- Google doc for lecture questions
 - See Piazza for link, section 102
 - https://docs.google.com/document/d/1G6hkekQS7mT9lFpP8AVftYao8vLRujIrRLAvOuX_07w/edit



- Add your question anonymously (at the top)
- Help answer questions too!

Overview

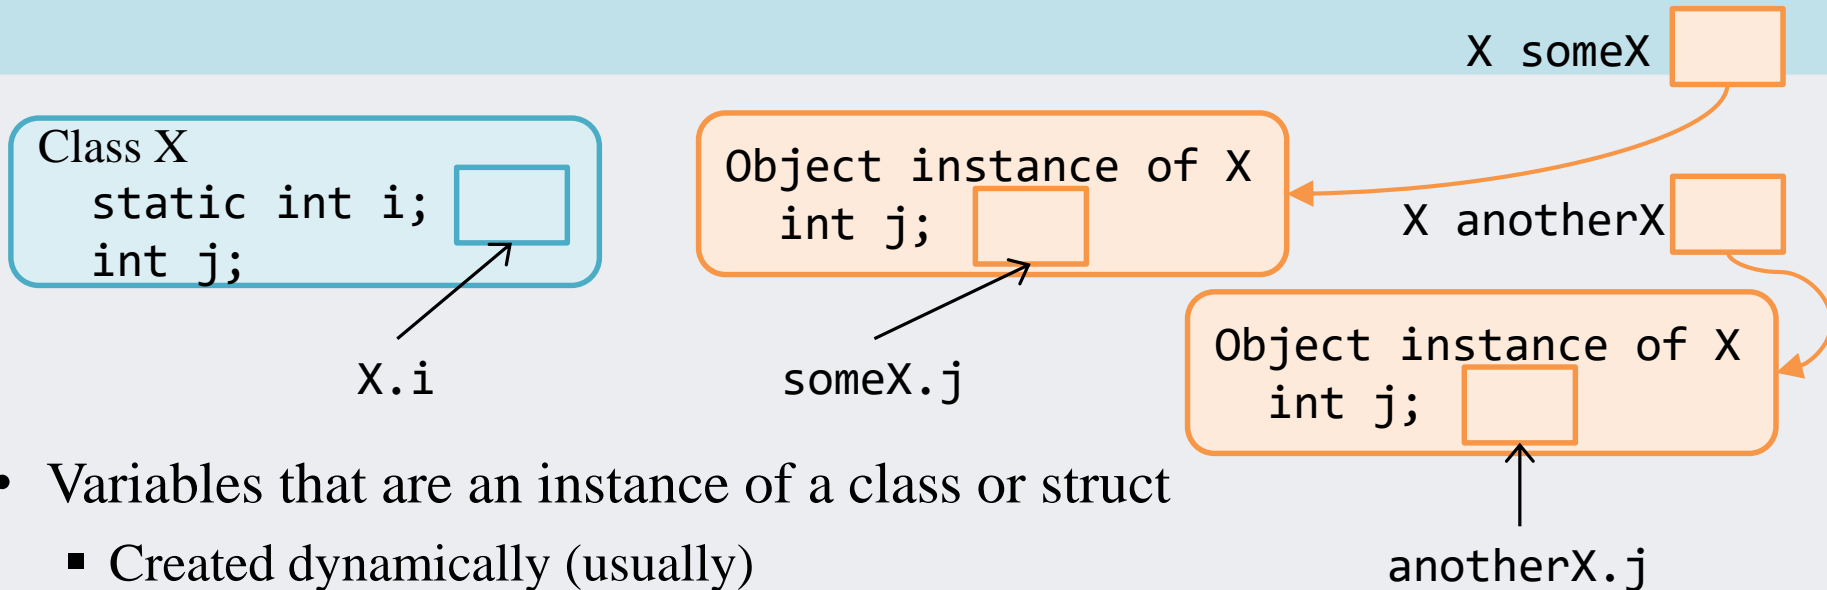
- Reading
 - Companion: 2.4.4-2.5
- Reference
 - Textbook: 3.9.1, 9.9, 3.10
- Learning objectives
 - read and write C code that includes structs
 - compare Java classes/objects with C structs
 - explain the difference between static and non-static variables in Java and C
 - explain why ISAs have both base-plus-offset and indexed addressing modes by showing how each is useful for implementing specific features of programming languages like C and Java
 - distinguish static and dynamic computation for access to members of a static struct variable in C
 - distinguish static and dynamic computation for access to members of a non-static struct variable in C
 - translate C struct-access code into assembly language
 - count memory references required to access struct elements
 - compare Java dynamic allocation the C's explicit-delete approach by identifying relative strengths and weaknesses
 - identify and correct dangling-pointer and memory leak bugs in C caused by improper use of free()
 - write C code that uses techniques that avoid dynamic allocation as a way to minimize memory-allocation bugs
 - write C code that uses reference counting as a way to minimize memory-allocation bugs
 - explain why Java code does not have dangling-reference errors but can have memory-leak errors

Variable types seen so far...

The simplest variables in any language

- Static variables
 - the compiler allocates them; i.e. their memory address is a constant
- Scalars and arrays
 - a scalar variable stores a single value
 - an array stores multiple values named by a single variable and an index
 - array data can be allocated either statically or dynamically
 - array variable can *be* the array (static) or store a pointer to it (dynamic)
 - index is generally a dynamic value

Instance variables



- Variables that are an instance of a class or struct
 - Created dynamically (usually)
 - many instances of the same class/struct can coexist
- Java vs C
 - Java: objects are instances of non-static variables of a class
 - C: structs are named variable groups, instance also called a struct
- Accessing an instance variable
 - requires a reference/pointer to a particular object/struct
 - then variable name chooses a variable in that object/struct

Structs in C

- A struct is a collection of variables of arbitrary type
 - allocated and accessed together

- Declaration


- similar to declaring a Java class without methods
- name is "struct" plus name provided by programmer

- static: `struct D d0;`
- dynamic: `struct D* d1;`

- Access

- static: `d0.e = d0.f;`
- dynamic: `d1->e = d1->f;`

$$d1 \rightarrow e \equiv (*d1).e$$



type

```
struct D {  
    int e;  
    int f;  
};
```

\approx

```
class D {  
    public int e;  
    public int f;  
};
```

can also access variables by dereferencing pointer and using "."

Struct allocation

```
struct D {  
    int e;  
    int f;  
};
```

- Static structs are allocated by the compiler

```
struct D d0;
```

Static memory layout

```
0x1000: value of d0.e  
0x1004: value of d0.f
```


- Dynamic structs are allocated at runtime
 - the variable that stores the struct pointer may be static or dynamic
 - the struct itself is allocated when the program calls `malloc`

```
struct D* d1;
```

Static memory layout

```
0x1000: value of d1 (address of struct)
```

```
struct A {  
    int x ; 4 bytes  
    int y ; 4 bytes  
    int* z ; 4 bytes  
};
```



Struct allocation

Dynamic allocation

```
struct D {  
    int e;  
    int f;  
};
```

- Runtime allocation of dynamic struct

(static) `struct D* d1;`

```
void foo() {  
    d1 = malloc( sizeof(struct D) );  
};
```

*.pos 0x1000
d1: .long 0*

- assuming that the code above allocates the struct at address 0x2000:

d1: 0x1000: 0x2000 (address of struct)
...
0x2000: value of d1->e
0x2004: value of d1->f

Struct access

```
struct D {  
    int e;  
    int f;  
};
```

in assembly

- Struct members can be accessed using **offset** from base address
 - offset to each member/variable from base of struct is **static**
- As before, static and dynamic differ by an extra memory access

```
d0.e = d0.f;
```

(d0 is static)

```
0x1000: value of d0.e  
0x1004: value of d0.f
```

```
m[0x1000] ← m[0x1004]
```

```
r[0] ← 0x1000
```

r0 = base address

```
r[1] ← m[r[0] + 4]  
m[r[0]] ← r[1]
```

*# r1 = d0.f
d0.e = r1*

??

```
d1->e = d1->f;
```

d1

```
0x1000: 0x2000
```

...

```
0x2000: value of d1->e
```

```
0x2004: value of d1->f
```

```
m[m[0x1000]] ← m[m[0x1000] + 4]
```

```
r[0] ← 0x1000
```

address of d1 pointer

```
r[1] ← m[r[0]]
```

r1: base address of struct

```
r[2] ← m[r[1] + 4]
```

```
m[r[1]] ← r[2]
```

```
struct D {
    int e;
    int f;
};
```

Struct access

in assembly

```
d0.e = d0.f;
```

```
0x1000: value of d0.e
0x1004: value of d0.f
```

```
m[0x1000] ← m[0x1004]
```

```
r[0] ← 0x1000
```

```
r[1] ← m[r[0] + 4]
m[r[0]] ← r[1]
```

```
ld $0x1000, r0 # r0 = address of d0
```

```
ld 4(r0), r1 # r1 = d0.f
st r1, (r0)  # d0.e = d0.f
```

```
d1->e = d1->f;
```

```
0x1000: 0x2000
```

...

```
0x2000: value of d1->e
0x2004: value of d1->f
```

```
m[m[0x1000]] ← m[m[0x1000] + 4]
```

```
r[0] ← 0x1000
```

```
r[1] ← m[r[0]]
```

```
r[2] ← m[r[1] + 4]
```

```
m[r[1]] ← r[2]
```

```
ld $0x1000, r0 # r0 = address of d1
```

```
ld (r0), r1 # r1 = address of struct
```

```
ld 4(r1), r2 # r2 = d1->f
```

```
st r2, (r1) # d1->e = d1->f
```

- Introducing... the load/store base **plus offset** instruction
 - dynamic base address in a register, plus a static offset (displacement)

```
ld 4(r1), r2
```

SM213 ISA

Revised to include offsets

- Machine format for base + offset:
 - offset in our case will always be a multiple of 4
 - we have only 4 bits to store the offset
 - so, we will store $p = o/4$ in the instruction

Name	Semantics	Assembly	Machine
load immediate	$r[d] \leftarrow v$	ld \$v, rd	0d-- vvvvvvvv
load base + offset	$r[d] \leftarrow m[r[s]+o]$	ld o(rs), rd	1psd
load indexed	$r[d] \leftarrow m[r[s]+4*r[i]]$	ld (rs, ri, 4), rd	2sid
store base + offset	$m[r[d]+o] \leftarrow r[s]$	st rs, o(rd)	3spd
store indexed	$m[r[d]+4*r[i]] \leftarrow r[s]$	st rs, (rd, ri, 4)	4sdi

Memory addressing modes reviewed

- Scalars

```
i = a;
```

- address in register (e.g. r1)
- access memory at address in register

```
ld (r1), r0
```

- Arrays

```
i = a[j];
```

- **base address** in register (e.g. r1 stores address of a)
- dynamic **index** in register (e.g. r2 stores value of j)
- access memory at base + index * element size (e.g. 4)

```
ld (r1, r2, 4), r0
```

- Struct members (instance variables)

```
i = a.j;
```

```
i = b->k;
```

- base address in register (e.g. r1 stores starting address of struct)
- static (constant) **offset** (e.g. X is offset to j/k from beginning of struct)
- access memory at base plus offset
- equivalent to array access with static index

```
ld X(r1), r0
```

Variations in struct declaration

- struct variables can be declared inside other structs



```
struct D {  
    int e;  
    int f;  
};
```

```
struct X {  
    int i;  
    struct D d;  
    int j;  
};
```



```
struct Y {  
    int i;  
    struct D* d;  
    int j;  
};
```

```
struct Z {  
    int i;  
    struct Z* z;  
    int j;  
};
```

struct D layout:



struct Y



- struct members can be arrays or pointers

W contains
the
actual
array

```
struct W {  
    int i;  
    int a[10];  
    int j;  
};
```



```
struct U {  
    int i;  
    int* a;  
    int j;  
};
```

array
lives
somewhere
else

- What is the offset (in bytes) of member `j` in `struct X` below?

```
struct D {  
    int e;  
    int f;  
};
```

```
struct X {  
    int i;  
    struct D d;  
    int j;  
};
```



- A. 0
- B. 4
- C. 8
- ☒ D. 12
- E. something else

- What is the offset (in bytes) of member `j` in `struct X` below?
Assume 4-byte addresses.

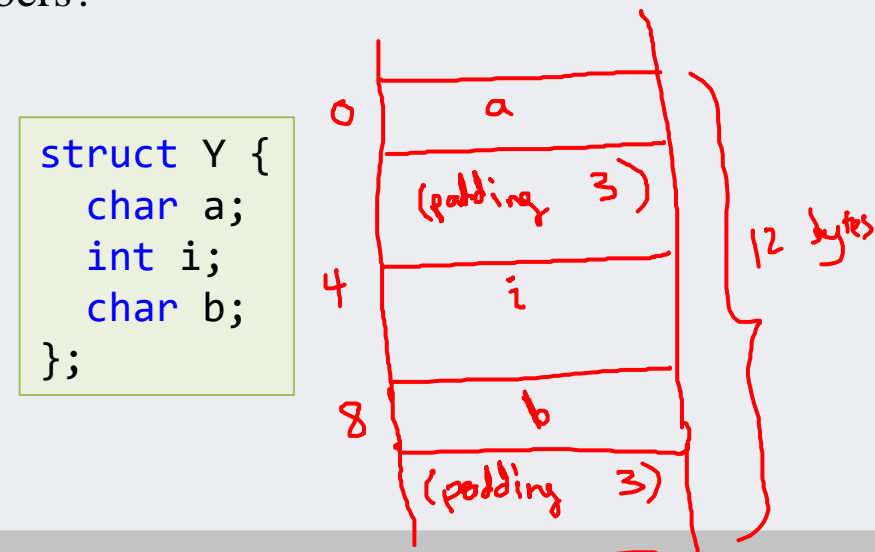

```
struct D {  
    int e;  
    int f;  
};
```

```
struct X {  
    0 int i;  
    4 struct D* d;  
    8 int j;  
};
```

- A. 0
- B. 4
- ☒ C. 8
- D. 12
- E. something else

Struct size and alignment

- Alignment rules apply inside of a struct
 - Each instance variable will be aligned within the struct according to its type size
 - structs are aligned according to their largest instance variable type
 - Structs can mix types (**padding** might be added in or after the members)
 - What are the *sizes* of the structs below? (i.e. `sizeof(struct X)`)
 - what are the *offsets* of each of their members?



Arrays of structs

- We can declare arrays of structs (or of struct pointers)

```
struct X {  
    int i;  
    int j;  
    int k;  
};
```

```
struct X a[10];  
struct X* b[10];  
  
...  
  
a[3].i = b[5]->j;
```

array
of
struct X
one
struct X

array
of
pointers
one
pointer

- What is the offset (in bytes) of member j in struct X below?

```
struct D {  
    int e; 4  
    int f; 4  
}; 8 bytes
```

```
struct X {  
    int i; 4  
    struct D d[10]; 80 bytes  
    int j;  
}; 84
```

Ans. 84

iClicker 1c.4


short: 2 bytes

- What is the offset (in bytes) of member j in struct X below?

```
struct Y {  
    char c;  
    short s;  
};
```

struct Y

```
struct X {  
    int i; 4  
    struct Y y[10]; 40  
    int j; 44  
};
```




Ans 44

- Given

44 bytes { struct A {
 char a; 1+3 padding
 int b[10];
};

```
struct X {
  char c;
  struct A a;
  int j;
};
```



```
struct X s; // a global variable
```

```
ld $s, r0
```

- Which of the following loads `s.a.b[2]` into `r1`?

- A. `ld 10(r0), r1`
- B. `ld 12(r0), r1`
- C. `ld 16(r0), r1`
- D. `ld 20(r0), r1`
- E. `ld 24(r0), r1`

- Given

```
struct A {  
    char a;  
    int b[10];  
    int* p;  
};
```

```
struct X {  
    char c;  
    struct A a;  
    int j;  
};
```

```
struct X s; // a global variable
```

```
ld $s, r0
```

- Which of the following loads `s.a.b[2]` into `r1`?

A. `ld 10(r0), r1`

B. `ld 12(r0), r1`

☒ C. `ld 16(r0), r1`

D. `ld 20(r0), r1`

E. `ld 24(r0), r1`

- Given


```
struct X {
    int a[2];
    char c;
};
```

Handwritten notes: 2×4 (for int array), 1 (for char), $+3$ padding

```
struct X s; // a global variable
```

- What is the value of `sizeof(s)`?

- A. 3
- B. 9
- ☒ C. 12
- D. 16



- Given

```
struct A {
    char a;
    4 int* b;
    int* p;
};
```

```
struct X {
    char c;
    struct A* a;
    int j;
};
```

```
struct X* s; // a global variable
```

ld \$s, r0 # r0: address of s pointer

ld (r0), r0 # r0: address of the struct X

ld 4(r0), r0 # r0: value of s → a
(address of struct A)

ld 4(r0), r0 # r0: value of s → a → b
(address of 'int array')

ld 8(r0), r1 # r1: s → a → b[2]

- How many memory reads to load s → a → b[2] into r1?


A. 2

B. 3

☒ C. 4

D. 5

E. 6



- Given


```
struct A {
    char a;
    int b[10];
    int* p;
};
```

3 bytes padding

```
struct X {
    char c;
    struct A a;
    int j;
};
```

3 bytes padding

```
struct X s; // a global variable
```



- How many memory reads to load `s.a.b[2]` into `r1`?

A. 1

B. 2

C. 3

D. 4

E. 5

`ld $s, r0 # r0: base address of s`

- `ld 16(r0), r1 # r1: s.a.b[2]`

- Given

```
struct A {
  char a;
  int* b;
  int* p;
};
```

3 bytes padding

```
struct X {
  char c;
  struct A a;
  int j;
};
```

3 bytes padding

```
struct X s; // a global variable
```

ld \$s, r0

- ld 8(r0), r0*
- ld 8(r0), r1*

- How many memory reads to load `s.a.b[2]` into `r1`?


A. 1

B. 2

C. 3

D. 4

E. 5




Dynamic allocation in C and Java

- Programs can allocate memory dynamically
 - allocation reserves a range of memory for a purpose
- In Java, instances of classes are allocated by the `new` keyword
- In C, byte ranges are allocated by call to `malloc` procedure
 - these bytes can be used for `any type` that can fit in them

Dynamic allocation in C

- Memory allocation: `void* malloc(unsigned long n);`

- n is the number of bytes to allocate
- return type is void*
 - a pointer to anything (no specific type assigned) "just an address"
 - can be cast to/from any other pointer type
 - cannot be dereferenced directly



- Use `sizeof` to determine the number of bytes to allocate
 - `sizeof(x)` statically computes the # of bytes in a type or variable
 - caution: `sizeof(pointer)` gives the size of a *pointer*, not what it *points to*

```
struct Foo* f = malloc(sizeof(*f));
```

```
void* p = malloc(sizeof(struct Foo));    *((int*) p) = 4;  
void* q = malloc(sizeof double);
```

possible, but dangerous!

Memory deallocation

- Wise management of memory requires **deallocation**
 - memory is a finite resource
 - deallocation frees previously allocated memory for re-use
- In Java:
 - automatic garbage collection: requires keeping track of every reference to an object
- In C:
 - Dynamic memory must be deallocated explicitly by calling **free**
 - **free** deallocates memory immediately; does not check to see if it is still in use
- **Memory leak:**
 - when dynamically allocated data is not deallocated when it is no longer needed
 - size of program gradually increases; available memory leaks away

Memory heap


- The **heap** is a large section of memory from which **malloc** allocates objects
 - **malloc** finds unused space in the heap, marks the chunk of bytes as used, and returns the address of the first byte
 - **free** marks the previously allocated chunk of bytes as unused
- In Java: all objects are stored on the heap

malloc (number of bytes) : returns address of first byte
free (address of an allocation)

Issues with explicit deallocation

must be an address associated with a malloc call

- What `free(x)` does
 - deallocates "object" at address `x` (`x` is a pointer)
 - this memory can be reused by subsequent call to `malloc`
- What `free(x)` does not do
 - it does not change the value of `x`
 - other variables may still point there too
 - the binary data stored at address `x` is not erased



Issues with explicit deallocation

```
struct buffer* create() {  
    struct buffer* buf = malloc(sizeof(*buf));  
    ...  
    return buf;  
}
```


malloc(sizeof(struct buffer));

```
void foo() {  
    struct buffer* mb = create();  
    bar(mb);  
    free(mb);  
}
```

```
struct buffer* aMB;
```

```
void bar(struct buffer* mb) {  
    aMB = mb;  
}
```

```
void bat() {  
    aMB->x = 0;  
}
```




- what might happen in `bar()`, and
- why a subsequent call to `bat()` would expose a serious bug?

Dangling pointers

"use after free"

- A dangling pointer is a pointer to an object that has been freed
 - or, could point to unallocated memory or to another object
- Why they are a problem
 - program thinks it is writing to object of type X, but isn't actually
 - it may be writing to an object of type Y



Avoiding dangling pointers in C

- Understand the problem
 - when allocation and free appear in different places in your code
 - e.g. when a procedure returns a pointer to something it allocates
- Avoid the problem cases, if possible
 - restrict dynamic allocation/free to single procedure
 - don't write procedures that return pointers
 - use local variables instead
 - we'll see later that local variables are automatically allocated on call and freed on return
- Engineer for memory management
 - define rules for which procedure is responsible for deallocation
 - implement explicit **reference counting** if multiple potential deallocators
 - define rules for which pointers can be stored in data structures
 - use coding conventions and documentation to ensure rules are followed

Co-locate allocation and deallocation

- If a procedure returns value of dynamically allocated object
 - allocate that object in its caller, and pass pointer to it to the procedure (callee)
 - good if caller does both `malloc` / `free` itself

```
struct buffer* receive() {  
    struct buffer* buf = malloc( sizeof(*buf) );  
    ...  
    return buf;  
}
```

```
void foo() {  
    struct buffer* mb = receive();  
    free(mb);  
}
```



```
void receive(struct buffer* buf) {  
    ...  
}
```

Delegate problem to caller

```
void foo() {  
    struct buffer* mb = malloc( sizeof(*mb) );  
    receive(mb);  
    free(mb);  
}
```

Transfer `malloc` to foo

Use local variables instead

- If a procedure does both malloc and free
 - use a "static" local variable instead of malloc
 - and don't give another procedure a pointer to the local variable
 - local variables are allocated on call and deallocated on return

```
struct buffer* receive() {  
    struct buffer* buf = malloc( sizeof(*buf) );  
    ...  
    return buf;  
}
```

```
void foo() {  
    struct buffer* mb = receive();  
    free(mb);  
}
```



```
void receive(struct buffer* buf) {  
    ...  
}
```

```
void foo() {  
    struct buffer mb;  
    receive(&mb);  
}
```


local struct buffer variable

automatically allocated on call,
passed by reference to receive

Example – C strings

- A string "like this" in C:

- is an array of `chars`
- has its end indicated by the first null `'\0'` in the array
- so, every string has a
 - maximum length (the capacity of the array – 1), and
 - a printable length determined by the position of the first null character



- The standard C library (in `string.h`) has many operations on strings
 - e.g. `strlen(s)` returns the (printable) length of a string
- Let's consider:
 - how to create a new string that is a copy of an existing string

*char**

String copy

Version 1

```
char* copy(char* s) {  
    int len = strlen(s);  
    char* d = malloc(len + 1);  
  
    for (int i = 0; i <= len; i++)  
        d[i] = s[i];  
  
    return d;  
}
```

C library equivalent: `strdup`

```
// in your application program  
void foo(char* s) {  
    char* d = copy(s);  
    printf("%s\n", d);  
  
    free(d);  
}  
  
void bar() {  
    foo("Hello, World!");  
}
```

- What can be improved in this code?

- delegate allocation, so that malloc and free are in the same procedure

String copy

Version 2

```
void copy(char* d, char* s) {  
    int len = strlen(s);  
    for (int i = 0; i <= len; i++)  
        d[i] = s[i];  
}
```

length for copying

```
// in your application program  
void foo(char* s) {  
    int len = strlen(s);  
    char* d = malloc(len + 1);  
    copy(d, s);  
    printf("%s\n", d);  
    free(d);  
}
```

length for allocation

C library equivalent: `strcpy`

- Pros and cons of this version?

what if d is much shorter than s?


String copy

Version 3

```
void copy(char* d, int dsize, char* s) {  
    int len = strlen(s);  
  
    if (len > dsize-1) len = dsize-1;  
  
    for (int i = 0; i < len; i++)  
        d[i] = s[i];  
  
    d[len] = '\0';  
}
```

C library equivalent: strncpy

```
// in your application program  
void foo(char* s) {  
    int len = strlen(s);  
    char* d = malloc(len + 1);  
    copy(d, len+1, s);  
    printf("%s\n", d);  
  
    free(d);  
}
```



"...if possible..."

... use local variables, don't return pointers, malloc/free in same procedure

- Sometimes, it's not possible
- If a reference needs to be passed among multiple modules
 - each module itself knows when it needs and doesn't need the reference
 - but these modules may not communicate that information between themselves
- What do we need to know in order to correctly **free** an object?
 - **free** should only happen if the object is still in use
 - you need to keep track of all its uses
 - when it starts being used
 - when it stops being used

Reference counting

- We can use **reference counting** to track object use
 - any procedure that stores a reference (starts using the object) increments the count
 - any procedure that discards a reference (stops using the object) decrements the count
 - the object is **freed** when the count goes to zero

Reference counting

One possible approach

- Reference counter can be part of a **struct**
 - updates to reference counter are done through special methods
 - no longer explicitly call **malloc/free** directly when creating the struct instances

```
struct buffer* malloc_buf() {  
    struct buffer* buf = malloc(sizeof *buf);  
    buf->ref_count = 1;  
    return buf;  
}
```

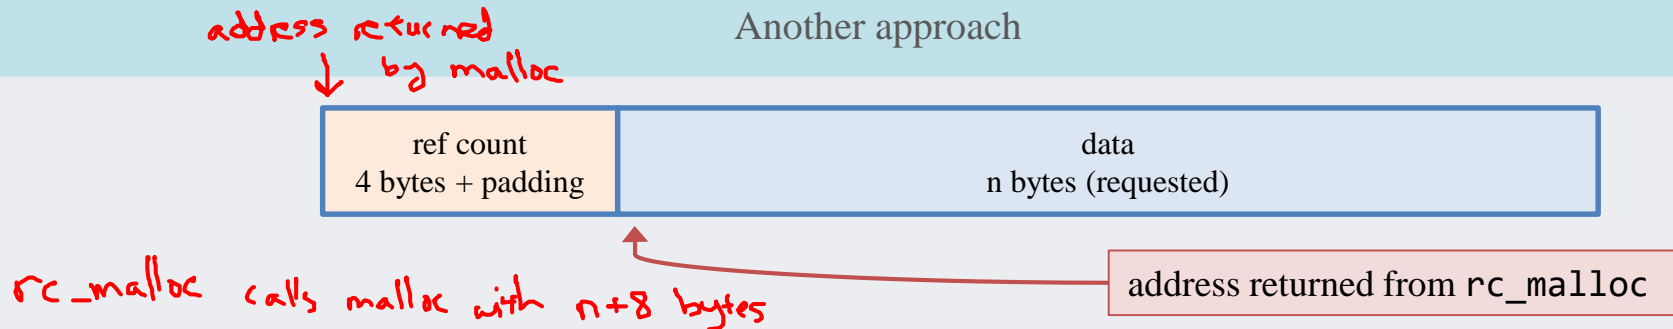
```
void add_reference(struct buffer* buf) {  
    buf->ref_count++;  
}
```

```
void free_reference(struct buffer* buf) {  
    buf->ref_count--;  
    if (buf->ref_count == 0)  
        free(buf);  
}
```

```
struct buffer {  
    ...  
    // the actual data attributes  
    ...  
    int ref_count;  
}
```

Reference counting

Another approach



- Sample code on Canvas: [ref-count-grades-example.zip](#)
 - refcount.c: malloc/free wrapper with ref-count implementation
 - map.c: hashmap where values are ref-counted
 - intlist.c: ref-counted list of integers
 - grades.c: implements a map from student ID to grade list
- Key memory management challenge
 - map.c and grades.c have references to grade lists
 - lists may be released in different points of the code
 - mapping may change from one value to another
 - marks may still be in use when mapping changes

Implementing reference counting

refcount.h

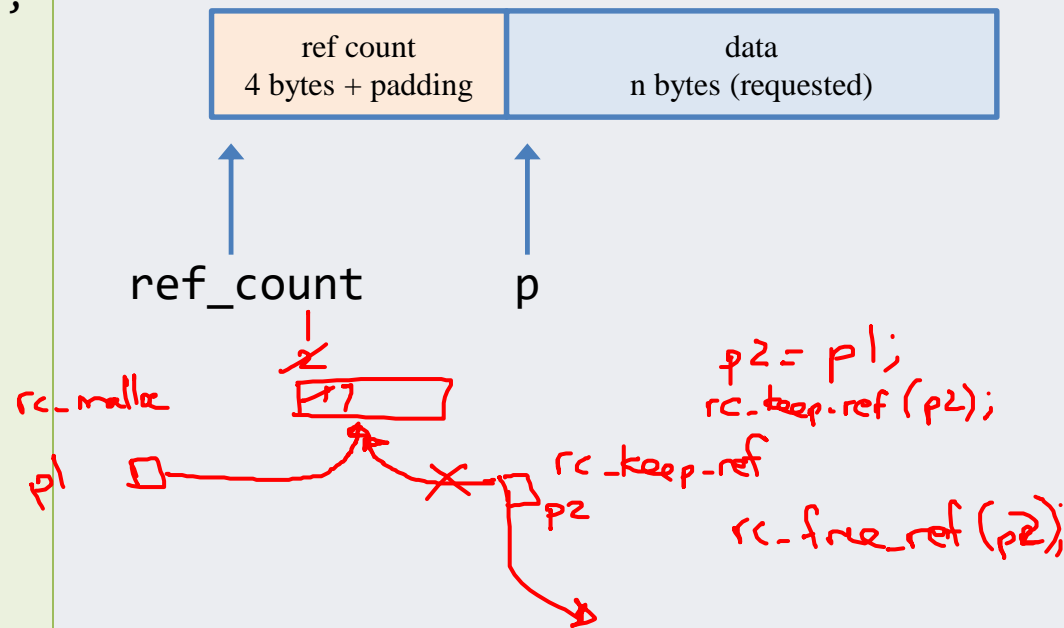
```
void* rc_malloc(int nbytes); // alternative malloc with room for ref_count
void rc_keep_ref(void* p);   // call when reference added
void rc_free_ref(void* p);   // call when reference removed
```

refcount.c

```
void* rc_malloc(int nbytes) {
    int* ref_count = malloc(nbytes + 8);
    *ref_count = 1;
    return ((void*) ref_count) + 8;
}


void rc_keep_ref(void* p) {
    int* ref_count = p - 8;
    (*ref_count)++;
}

void rc_free_ref(void* p) {
    int* ref_count = p - 8;
    (*ref_count)--;
    if (*ref_count == 0)
        free(ref_count);
}
```



Reference counting demonstration


stack.zip, on Canvas



Reference counting and procedure calls

refcount_cases.c, on Canvas


- When a reference is a parameter
 - the caller has a reference and maintains it for the duration of the call
 - so, the callee need not perform a `keep_ref` UNLESS
 - callee saves the pointer someplace that will outlast the call (e.g. global variable)
 - callee returns a pointer to the caller



Reference counting and procedure calls

refcount_cases.c, on Canvas


- When a reference is a return value
 - the callee must have a reference to the value
 - it passes that reference to the caller
 - the callee implicitly gives up its reference as part of the return
 - the reference is transferred to the caller
 - the caller must call `free_ref` when it no longer stores the reference



Reference counting and procedure calls

refcount_cases.c, on Canvas

- When a reference is stored in a local variable
 - that variable goes away implicitly when the procedure returns
 - and so the procedure must call `free_ref` before it returns



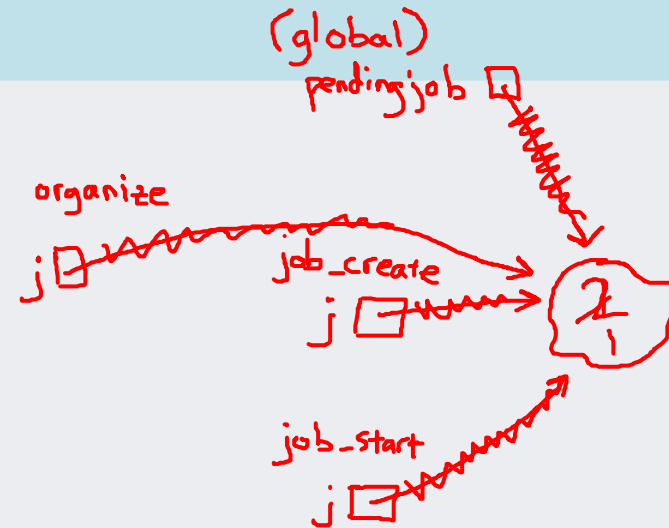
OK, let's make this an iClicker (four of them!)

job.zip on Canvas

- Select ALL of the following that are true.
- Line numbers refer to the original program

1. → `job.c`, between lines 26 and 27 **B**
2. → `job.c`, between lines 32 and 33 **A**
3. → `job.c`, between lines 35 and 36 **D**
4. → Replace in `organize.c`, line 8 **C**

- A. add `rc_keep_ref`
- B. add `rc_free_ref`
- C. replace line(s)
- D. no change




Be careful with references

quiz.zip, on Canvas

- Consider this code:

```
int main(int argc, char** argv) {
    char* name = strdup("Quiz X");
    quiz_t quizzes[4];
    for (int i = 0; i < 4; i++) {
        name[5] = i + 49; // ASCII numeral
        quizzes[i] = quiz_create(name);
    }
    for (int i = 0; i < 4; i++) {
        printf("%d %s\n", i, quiz_get_name(quizzes[i]));
    }
}

quiz_t quiz_create(char* name) {
    struct quiz* quiz = malloc(sizeof(*quiz));
    quiz->name = name;
    return quiz;
}
```



It prints:

```
0 Quiz 4
1 Quiz 4
2 Quiz 4
3 Quiz 4
```

Why?


Nesting and reference counting

quiz.zip, continued

- How does this impact reference counting?

```
quiz_t quiz_create(char* name) {  
    struct quiz* quiz = malloc(sizeof(*quiz));  
    quiz->name = strdup(name);  
    return quiz;  
}  
  
void quiz_delete(quiz_t quiz) {  
    free(quiz->name);  
    free(quiz);  
}
```

What must change to use reference counting?



Nesting and reference counting

What if there is more to the quiz structure?


```
struct quiz {  
    char* name;  
    int* grades;  
    int num_grades;  
};
```

```
quiz_t quiz_create(char* name) {  
    quiz_t quiz = rc_malloc(sizeof(*quiz));  
    int name_size = strlen(name) + 1;  
    quiz->name = rc_malloc(name_size);  
    strcpy(quiz->name, name, name_size);  
    quiz->grades = NULL;  
    return quiz;  
}
```

```
void addGrades(quiz_t quiz, int* grades, int numGrades) {  
    quiz->grades = rc_malloc(sizeof(*quiz->grades) * numGrades);  
    for (int i = 0; i < numGrades; i++)  
        quiz->grades[i] = grades[i];  
}
```

```
void quiz_delete(quiz_t quiz) {  
    rc_free_ref(quiz->name);  
    rc_free_ref(quiz);  
}
```

How should we fix
quiz_delete?



Detecting memory problems

Valgrind

- Valgrind is a program that performs dynamic analysis of the runtime of a program
 - e.g. run using `valgrind ./my_program`
- It runs your program and monitors dynamic allocation and deallocation
- It can tell if your program has:
 - memory leaks
 - use after free (dangling pointers)

Garbage collection

Deallocation in Java

- In Java, objects are deallocated implicitly
 - the program never says **free**
 - the runtime system tracks **every object reference**
 - a garbage collector runs periodically to deallocate unreachable objects
- Advantage compared to explicit **free** (C)
 - No dangling pointers
 - (almost) no memory leaks
 - reference cycles are not a problem
 - e.g. object a has a pointer to b, object b has a pointer to a
 - these would cause memory leak when using reference counting alone



Discussion

- What are the advantages of explicit free?
fast, efficient, makes us careful with memory management
- What are the advantages of reference counting?
reduces inter-module coupling
- What are the advantages of garbage collection?
automatic, handles most cases
- Should we ignore deallocation in Java programs?
NO e.g. HashMap

Memory management in Java

- Memory "leaks" can still occur
 - when garbage collector fails to reclaim **unneeded** objects
 - still a significant problem for long-running programs where garbage accumulates
- Why would an object not be reclaimed?
 - garbage collector only reclaims **unreachable** objects
 - Unneeded/unused objects that keep references are not reclaimed
 - Collections and maps may maintain old unneeded objects