



**a place of mind**

THE UNIVERSITY OF BRITISH COLUMBIA

# CPSC 213

## Introduction to Computer Systems

### Unit 2a

#### I/O Devices, Interrupts, and DMA

All slides adapted from materials by Mike Feeley, Jonatan Schroeder, Robert Xiao, and Jordon Johnson

# Announcements

- Google doc for lecture questions
  - See Piazza for link, section 102
  - [https://docs.google.com/document/d/1G6hkekQS7mT9lFpP8AVftYao8vLRujIrRLAvOuX\\_07w/edit](https://docs.google.com/document/d/1G6hkekQS7mT9lFpP8AVftYao8vLRujIrRLAvOuX_07w/edit)



- Add your question anonymously (at the top)
- Help answer questions too!

# Overview


- Reading in Text

- 8.1, 8.2.1, 8.5.1-8.5.3

- Learning Goals

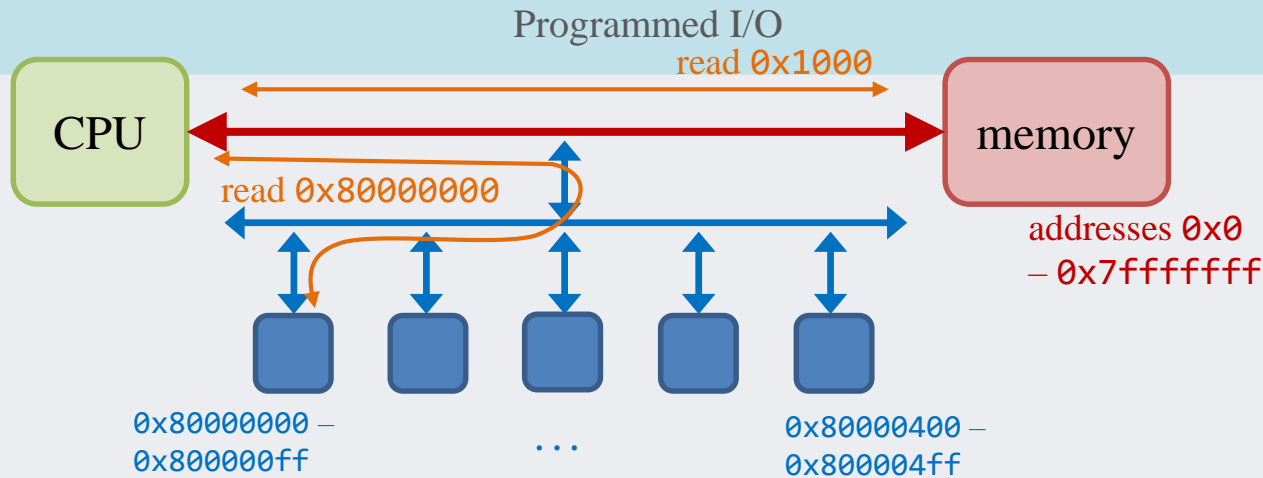
- Explain what PIO is, why it exists, what processor originates it, and how it is used
- Explain what DMA is, why it exists, what it does, and what processor originates it
- Explain what an interrupt is, why it exists, what it does, and what processor originates it
- Compare PIO and DMA by identifying when each *can* be used and when each *should* be used
- Compare PIO and interrupts by identify when each can be used and when each should be used
- Explain the relationship between polling, PIO and interrupts
- Explain the conditions that make polling acceptable or not
- Explain what happens when an interrupt occurs at the hardware level
- Explain what is means for an operation such as a disk read to be asynchronous and give examples of code that works when an operation is synchronous but does not work with it is asynchronous
- Write event-driven programs in C using function pointers
- Describe why event-driven programs may be harder to write, read, and debug

# Looking beyond the CPU and memory



- **Memory bus**
  - data/control path connecting CPU, main memory, and I/O bus
  - also called the *front-side bus*
- **I/O bus**
  - data/control path connecting memory bus and I/O controllers
  - e.g., PCI
- **I/O controller**
  - a processor running software (firmware)
  - connects I/O device to I/O bus
  - e.g., SCSI, SATA, Ethernet, ...
- **I/O device**
  - I/O mechanism that generates or consumes data
  - e.g., disk, radio, keyboard, mouse, ...

# Talking to an I/O controller



Real  
memory

I/O  
memory

- **Programmed I/O (PIO)**
  - CPU transfers a **word** <sup>→ unit of data</sup> at a time between CPU and I/O controller
  - typically use standard load/store instructions but to **I/O-mapped memory address**
- **I/O-mapped memory**
  - memory addresses beyond the end of main memory
  - used to name I/O controllers (usually configured at boot time)
  - loads and stores are translated into I/O bus messages to controller
- **Example:**
  - to read/write to controller at address 0x80000000:

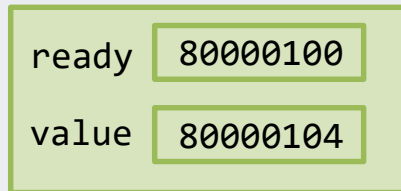
```
ld $0x80000000, r0
st r1, (r0)      # write the value of r1 to the device
ld (r0), r1      # read a word from the device into r1
```

# Limitations of PIO

- Reading or writing large amounts of data slows CPU
  - CPU must transfer one word at a time
  - controller/device is much slower than CPU
  - so, CPU runs at controller/device speed, mostly waiting for controller
- I/O controller cannot initiate communication
  - sometimes the CPU asks for data
  - but, sometimes the controller receives data for the CPU, without CPU asking
    - e.g. mouse click or network packet reception
  - how does controller notify CPU that it has data the CPU should want?
- One not-so-good idea:
  - what is it? *polling*
  - what are the drawbacks? *still have using PIO*
  - when is it OK?


# Polling and I/O memory

CPU



```
int pollDeviceForValue() {  
    volatile int* ready = 0x80000100;  
    volatile int* value = 0x80000104;  
  
    while (*ready == 0) {};  
  
    return *value;  
}
```


I/O controller !



```
void send(int v) {  
    volatile int* ready = 0x80000100;  
    volatile int* value = 0x80000104;  
  
    *value = v;  
    *ready = 1;  
}  
  
send(7);
```

- Reading (or writing) I/O memory is SLOW
  - CPU -> I/O device: "give me value at address x"
  - I/O device -> CPU: "OK, here is the value"
- Polling is OK if poll has low overhead, or has high probability of "hit"

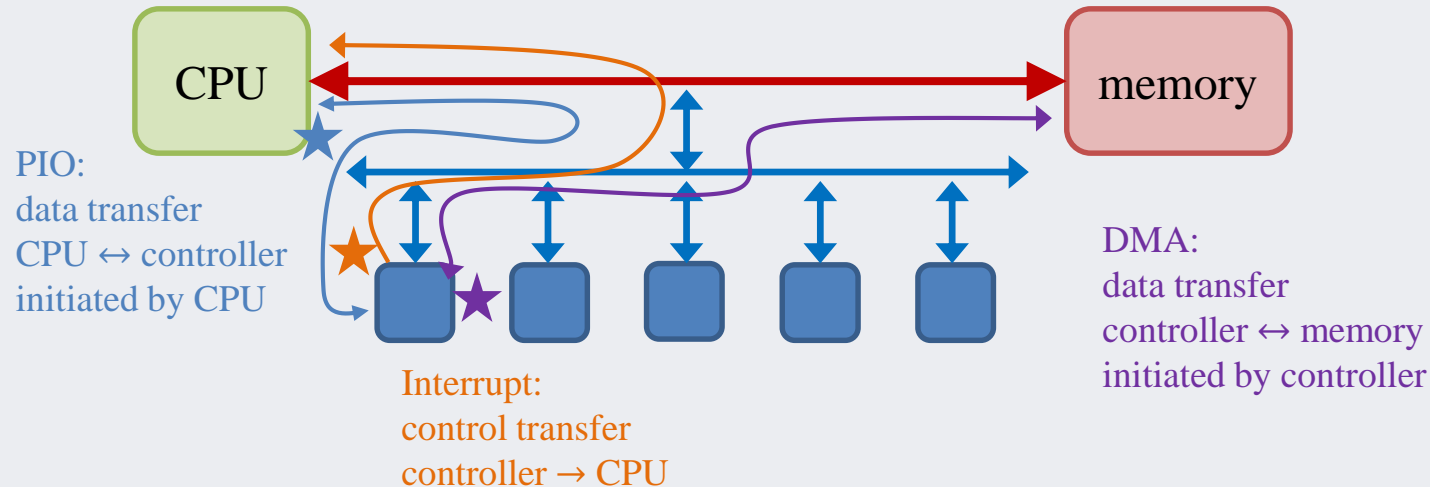
# Key observation



- CPU and I/O controller are independent processors
  - they should be permitted to work in parallel
  - either should be able to initiate data transfer to/from memory
  - either should be able to signal the other to get the other's attention



# Autonomous controller operation



- Direct Memory Access (DMA)

- controller can send/read data from/to any main memory address
- the CPU is oblivious to these transfers
- DMA addresses and sizes are *programmed* by CPU using PIO

- CPU Interrupts

- controller can signal the CPU
- CPU checks for interrupts on every cycle (its like a really fast, clock-speed poll)
- CPU jumps to controller's *Interrupt Service Routine* if it is interrupting

# Adding interrupts to SM213

- New special-purpose CPU registers
  - `isDeviceInterrupting` set by I/O controller to signal interrupt
  - `interruptControllerID` set by I/O controller to identify interrupting device
  - `interruptVectorBase` **interrupt-handler** jump table, initialized at boot time
- Modified fetch-execute cycle

```
while (true) {  
  if (isDeviceInterrupting) {  
    r[5] ← r[5] - 4; m[r[5]] ← r[6];  
    r[6] ← pc;  
    pc ← interruptVectorBase[interruptControllerID];  
  }  
  fetch();  
  execute();  
}
```

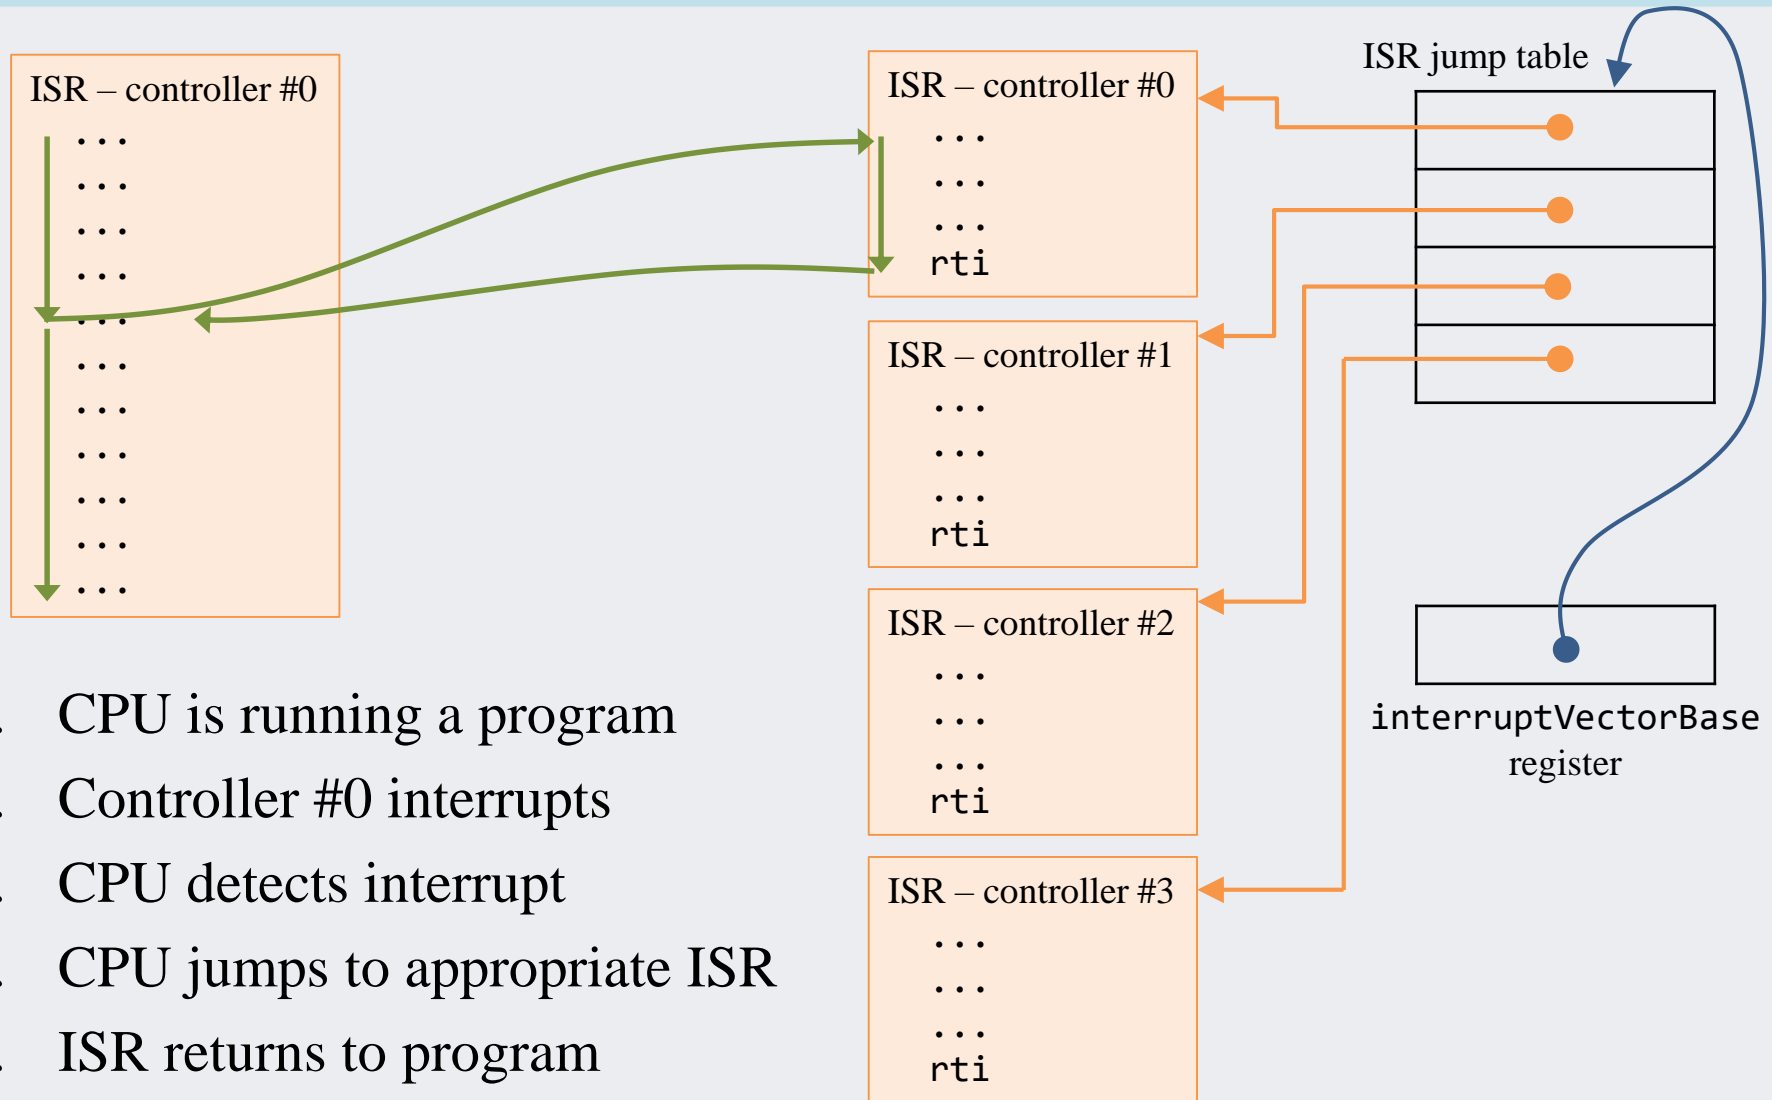
*'jump table' index*

**RTI: Return from Interrupt Instruction**

```
t ← r[6];  
r[6] ← m[r[5]]; r[5] ← r[5] + 4;  
isDeviceInterrupting ← 0;  
pc ← t;
```

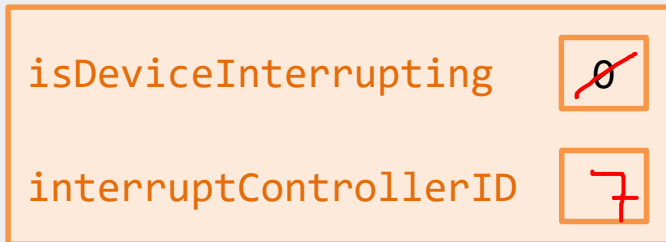
"polls" on CPU register instead of I/O memory

# Interrupt control flow on CPU



# Architectural view of interrupts

CPU



I/O controller #3

myID 3

I/O controller #7

myID 7

```
while (true) {  
  if (isDeviceInterrupting) {  
    r[5] ← r[5] - 4; m[r[5]] ← r[6];  
    r[6] ← pc;  
    pc ← interruptVectorBase[interruptControllerID];  
    isDeviceInterrupting ← 0;  
  }  
  fetch();  
  execute();  
}
```



**a place of mind**

THE UNIVERSITY OF BRITISH COLUMBIA

# Programming with I/O

# Disk read timeline

## CPU

1. PIO to request read

~~...  
idle waiting  
...~~    ...  
do other things  
...

6. Interrupt received  
Call `readComplete`

## I/O controller

2. PIO received, start read

...  
wait for read to complete  
...

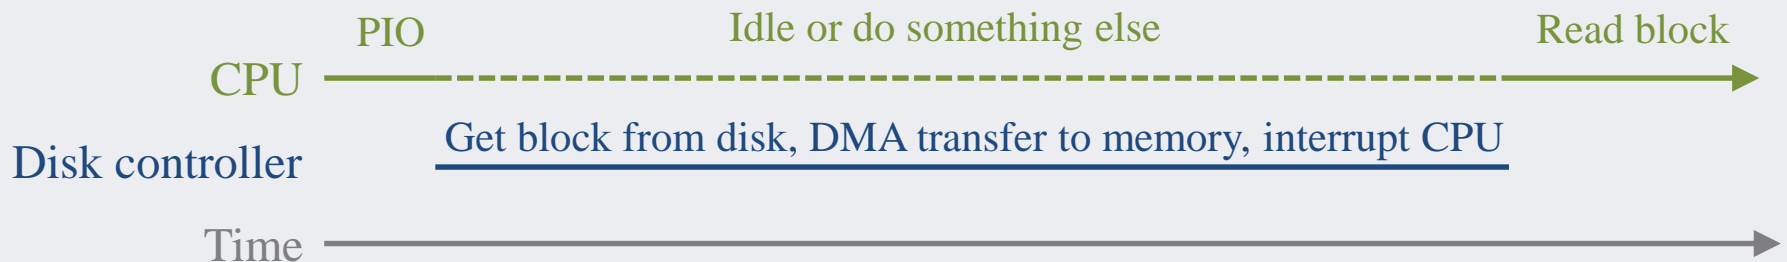
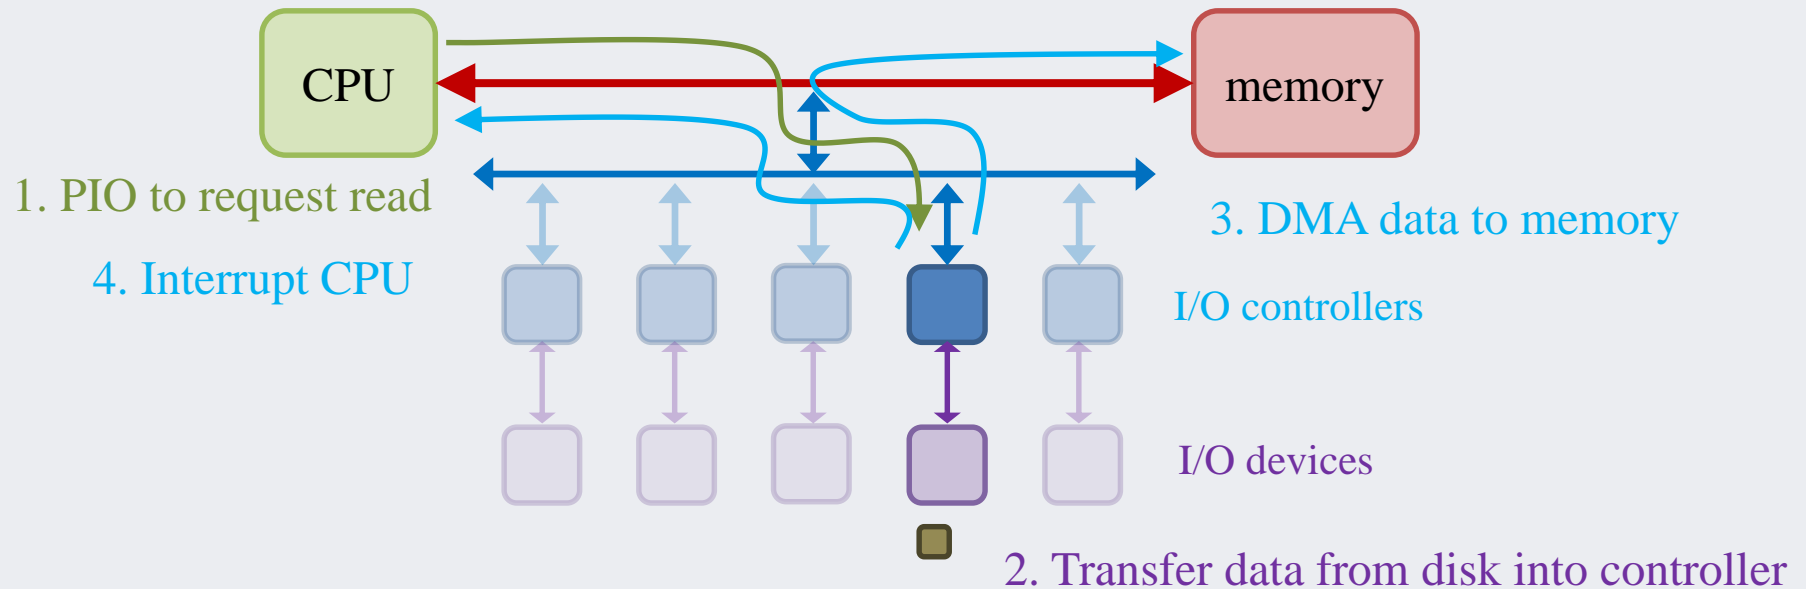
3. Read completes

4. Transfer data to memory (DMA)

5. Interrupt CPU

A disk stores blocks of data. Each block has a number.  
CPU can request read or write to a block

# Disk read timeline



# Sequential execution...

- Consider a program that reads data from a disk
  - you can think of the read as having three parts:
    1. figure out what data you want
    2. get the data
    3. do something with the data you got
  - these steps must happen in this order
  - we think of these steps as executing in sequence

- Example:

```
int sumDiskData (int blockNum, int numBytes) {  
    char buf[numBytes];  
    int sum = 0;  
  
    diskRead (buf, blockNum, numBytes);  
    for (int i = 0; i < numBytes; i++)  
        sum += buf[i];  
  
    return sum;  
}
```

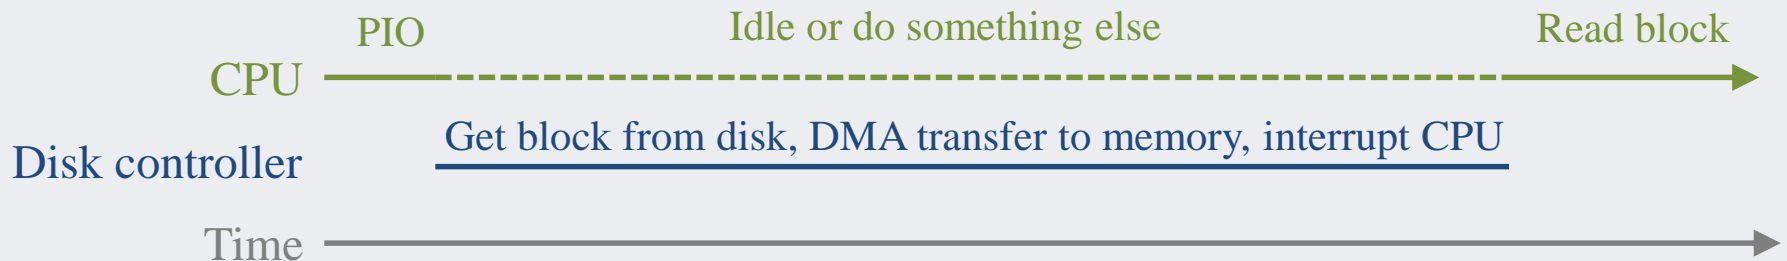


# ...meets reality

iClicker 2a.1

```
int sumDiskData (int blockNum, int numBytes) {  
    char buf[numBytes];  
    int sum = 0;  
  
    diskRead (buf, blockNum, numBytes);  
    for (int i = 0; i < numBytes; i++)  
        sum += buf[i];  
  
    return sum;  
}
```

Place these events on the timeline below:  
X: request `diskRead`  
Y: requested data is in `buf`  
Z: `for (int...)`



What is the actual order?

A: XYZ   B: XZY   C: Either XYZ or XZY, you can't tell   D: Some other order



# Making code asynchronous

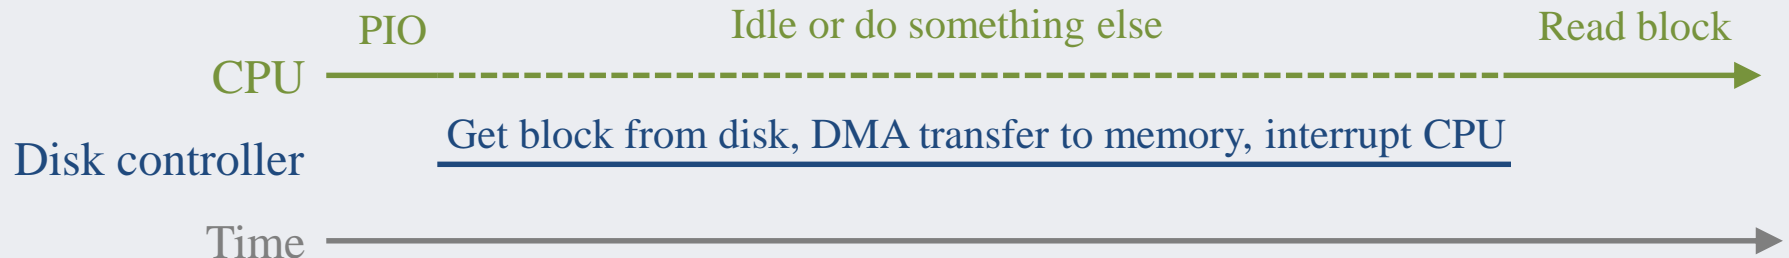
```
char buf[numBytes];  
diskRead (buf, blockNum, numBytes);
```

This code  
requests  
something

The desired execution has something happening  
between these two pieces of code

```
int sum = 0;  
for (int i = 0; i < numBytes; i++)  
    sum += buf[i];
```

This code should run when  
the request is complete



# Writing asynchronous code in C

- Events and event handlers

- things that cause asynchronous code to run are called **events**
  - e.g. disk-read completion
- the code that runs when an event occurs is called an **event handler** (or **callback**)
  - e.g. code that computes the checksum
- handlers are *registered* to a specific event
- events are *fired* to trigger the execution of the handler

- In code:

- request and registration of event handler are listed together in code
- this code continues, does not wait for event
- handler runs asynchronously when event occurs

```
void computeSum (char* buf, int blockNum, int numBytes) {
    int sum = 0;
    for (int i = 0; i < numBytes; i++)
        sum += buf[i];
}

int sumDiskBlock (blockNum, numBytes) {
    char buf [numBytes];

    diskRead (buf, blockNum, numBytes, computeSum);
}
```

# Making code asynchronous

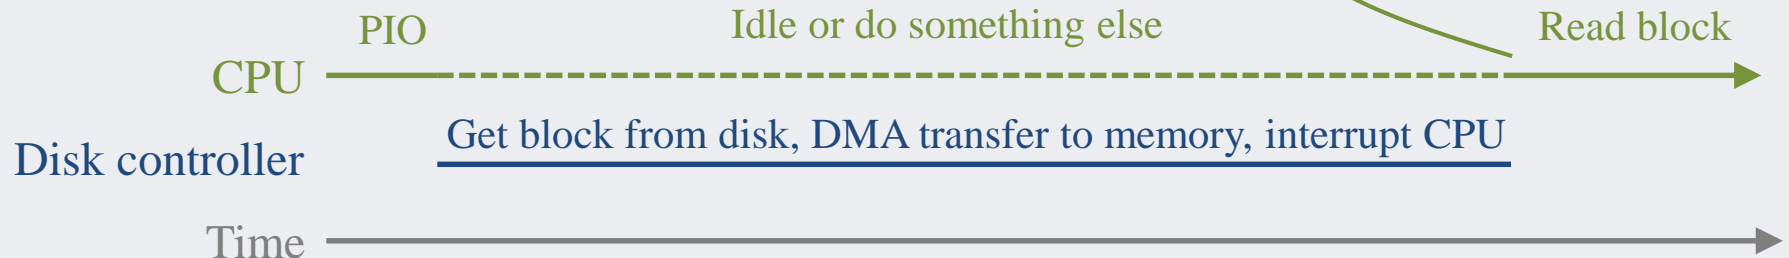
```
char buf[numBytes];  
diskRead (buf, blockNum, numBytes, computeSum);
```

Registration of handler

Event handler

```
void computeSum (char* buf, int blockNum, int numBytes) {  
    int sum = 0;  
    for (int i = 0; i < numBytes; i++)  
        sum += buf[i];  
}
```

Event fired



# Disk read PIO

- Expanded disk read with PIO
  - the message sent to disk has several fields
  - each field has different device-memory address
  - access it like a struct
    - but, keep in mind that writes are to device memory;
    - they are messages across bus to device controller; they are not writes to main memory

```
void* diskAddress = (void*) 0x80001000;
```

```
#DEFINE DISK_OP_READ 1
```

```
#DEFINE DISK_OP_WRITE 2
```

```
struct disk_ctl {  
    int    op;  
    char*  buf;  
    int    blockNum;  
    int    numBytes;  
};
```

```
void diskRead (char* buf, int blockNum, int numBytes, ...) {  
    struct disk_ctl* dc = diskAddress;  
    enqueue_handler (whenComplete, buf, blockNum, numBytes);  
    dc->op           = DISK_OP_READ;  
    dc->buf          = buf;  
    dc->blockNum     = blockNum;  
    dc->numBytes     = numBytes;  
}
```

# Implementing disk read (simplified)

- Interrupt vector, device ID, and PIO address
  - initialized before all this starts
  - by operating system when device connects
- Disk read
  - register event handler
  - request block (PIO)
- Interrupt handler
  - find event handler and fire event
  - firing means calling the handler procedure

```
#define MAX_DEVICES
void (*interruptVector [MAX_DEVICES])();

int diskID = 4;
int* diskAddress = (int*) 0x80001000;

interruptVector [diskID] = diskISR;
```

```
void diskRead (char* buf, int blockNum, int numBytes
               void (*whenComplete) (char*, int, int)) {
    enqueue_handler (whenComplete, buf, blockNum, numBytes);
    // perform PIO ... more about this later
}
```

```
void diskISR() {
    struct handler_dsc {
        void (*handler) (char*, int, int);
        char* buf;
        int blockNum;
        int numBytes;
    };
    struct handler_dsc hd;
    dequeue_handler (&hd);
    hd.handler(hd.buf, hd.blockNum, hd.numBytes);
}
```

# Did we really solve the problem?

iClicker 2a.2

- We wanted to do this:

```
int sumDiskBlock (int blockNum, int numBytes) {  
    char buf [numBytes];  
    int sum = 0;  
  
    diskRead (buf, blockNum, numBytes);  
    for (int i = 0; i < numBytes; i++)  
        sum += buf[i];  
  
    return sum;  
}
```

- But reality forced us to do this:

Which line is wrong?

A `int` computeSum (`char*` buf, `int` blockNum, `int` numBytes) {  
B `int` sum = 0;  
C `for` (`int` i = 0; i < numBytes; i++)  
D `sum` += buf[i];  
→ `return` sum;  
}

`void` sumDiskBlock (blockNum, numBytes) {  
 `char` buf [numBytes];  
  
 diskRead (buf, blockNum, numBytes, computeSum);  
}

What's the problem?

# Why can't computeSum return a value?

iClicker 2a.3

```
int computeSum (char* buf, int blockNum, int numBytes) {  
    int sum = 0;  
    for (int i = 0; i < numBytes; i++)  
        sum += buf[i];  
    return sum;  
}  
  
void sumDiskBlock (blockNum, numBytes) {  
    char buf [numBytes];  
  
    diskRead (buf, blockNum, numBytes, computeSum);  
}
```

What procedure calls computeSum?

- A. sumDiskBlock
- B. the procedure that calls sumDiskBlock
- C. another procedure in this program
- ☒ D. the disk interrupt service routine
- E. something else / I don't know



# Connecting asynchrony to program

- How do we use the value computed from the disk block?

- suppose we want to print it

```
void something () {  
    ...  
    int s = sumDiskBlock (blk, n);  
    printf("%d\n", s);  
}
```

- but asynchronously?

```
void computeSum (char* buf, int blockNum, int numBytes) {  
    int sum = 0;  
    for (int i = 0; i < numBytes; i++)  
        sum += buf [i];  
    free (buf);  
}  
void sumDiskBlock (blockNum, numBytes) {  
    char* buf = malloc (numBytes);  
    diskRead (buf, blockNum, numBytes, computeSum);  
}
```

# Ordering in asynchronous programs

- If something has to happen after an event
  - it must be triggered by the event
  - often this means it must be part of (or called by) event's handler
- To print the checksum

```
void computeSumAndPrint (char* buf, int blockNum, int numBytes) {  
    int sum = 0;  
    for (int i = 0; i < numBytes; i++)  
        sum += buf [i];  
    printf("%d\n", sum);  
    free (buf);  
}
```

Huge problem:

- often there is a ton of stuff that depend on the returned data
- these must happen after a particular event

```
void sumDiskBlock (blockNum, numBytes, whenComplete) {  
    char* buf = malloc (numBytes);  
    diskRead (buf, blockNum, numBytes, whenComplete);  
}  
  
void something() {  
    sumDiskBlock (1234, 4096, computeSumAndPrint);  
}
```

# Improving the code...

```
void computeSumAndPrint (char* buf, int blockNum, int numBytes) {  
    int sum = 0;  
    for (int i = 0; i < numBytes; i++)  
        sum += buf[i];  
    printf("%d\n", sum);  
    free(buf);  
}
```

← too specific, make this a parameter (callback)

```
void computeSumAndCallback (char* buf, int blockNum, int numBytes,  
                           void (*callback) (int)) {  
    int sum = 0;  
    for (int i = 0; i < numBytes; i++)  
        sum += buf[i];  
    callback(sum);  
    free(buf);  
}
```

## ...but making it worse

```
void computeSumAndCallback (... , void (*sumCallback) (int)) {
    int sum = 0;
    for (int i = 0; i < numBytes; i++)
        sum += buf[i];
    sumCallback(sum);
    free(buf);
}

void sumDiskBlock (blockNum, numBytes, whenComplete, sumCallback) {
    char* buf = malloc (numBytes);

    diskRead (buf, blockNum, numBytes, whenComplete, sumCallback);
}
```

```
void printInt (int i) {
    printf ("%d\n", i);
}

void something() {
    sumDiskBlock(1234, 4096, computeSumAndCallback, printInt);
}
```

What if you want to do something after printInt?

Welcome to "callback hell"...

# Happy system, sad programmer

- Humans like synchrony
  - we expect each step of a program to complete before the next one starts
  - we use the result of previous steps as input to subsequent steps
  - with disks, for example,
    - we read from a file in one step and then usually use the data we've read in the next step
- Computer systems are asynchronous
  - the disk controller takes 10-20 milliseconds ( $10^{-3}$ s) to read a block
    - CPU can execute 60 million instructions while waiting for the disk to complete one read
    - we must allow the CPU to do other work while waiting for I/O completion
  - many devices send unsolicited data at unpredictable times
    - e.g., incoming network packets, mouse clicks, keyboard-key presses
    - we must allow programs to be interrupted many, many times a second to handle these things
- Asynchrony makes programmers sad
  - it makes programs more difficult to write and much more difficult to debug

# Possible solutions

- Accept the inevitable
  - use an event-driven programming model
    - event triggering and handling are de-coupled
  - a common idiom in many Java programs
    - GUI programming follows this model
  - *CSP* is a language boosts this idea to first-class status
    - no procedures or procedure calls
    - program code is decomposed into a set of sequential/synchronous processes
    - processes can fire events, which can cause other processes to run in parallel
    - each process has a guard predicate that lists events that will cause it to run
  - Javascript in web browsers and Node.js embrace asynchrony, albeit awkwardly
- Invent a new abstraction
  - an abstraction that provides programs the illusion of synchrony
  - but, what happens when
    - a program does something asynchronous, like disk read?
    - an unanticipated device event occurs?
- What's the right solution?
  - we still don't know — this is one of the most pressing questions we currently face